

Operating systems

- Process synchronization
- Deadlocks

Process Synchronization

- Background
- Critical Section Problem
- Semaphores
- Classical Problems

Background

- A **cooperating** process is one that can affect or be affected by the other processes executing in the system.
- Any process that does **share** data with other processes is a **cooperating** process.
- **Concurrent access** to shared data may result in data **inconsistency**.
- Maintaining data consistency requires **mechanisms** to ensure the **orderly execution** of cooperating processes.

Cooperating processes

- Consider the following pseudo code:

```
global x  
x = 0  
x = x + 1  
print x
```

Value of x:
0
1
1<== output

Cooperating processes

- Two processes executing same code:

P0	P1	value of x:
global x	global x	
x = 0		0
<hr/>		
x = x + 1		1
<hr/>		
	x = 0	0
<hr/>		
print x		0 ← P0 output
<hr/>		
	x = x + 1	1
<hr/>		
	print x	1 ← P1 output

Critical Section Problem

- Each process has a segment of code, called **a critical section**, in which the process may be changing common variables, updating a table, writing a file, etc.
- When one process is executing in its critical section, **no other process** is to be allowed to execute in its critical section.
- Thus, the execution of critical sections by processes should be **mutually exclusive** in time.

Critical Section Problem

- Design a **protocol** which processes can use to cooperate.
- Each process must request **permission** to enter its critical section implemented in the **entry section**.

General structure of
a typical process



```
do {  
    entry section  
    critical section  
    exit section  
    remainder section  
} while (TRUE);
```

Critical Section Problem

- A solution to the critical-section problem must satisfy the following requirements:
 - ① **Mutual exclusion:** if process P_i is executing in its critical-section, then no other process can be executing in its critical-section.
 - ② **Progress:** If no process is in its critical-section, one of the processes that wish to enter their critical-sections will be selected to enter its critical-section.
 - ③ **Bounded waiting:** A bound on the number of selections that a process can wait for entering its critical-section.

Two-process solution

Process P₀

do {

while (**turn** != 0) {}

critical section

turn = 1

remainder section

} while (TRUE)

Process P₁

do {

while (**turn** != 1){}

critical section

turn = 0

remainder section

} while (TRUE)

turn is shared variable. Initially **turn** = 0 or 1. This algorithm satisfies **mutual exclusion**, but not **progress**.

No progress solution

- The problem is that algorithm doesn't retain information about the **state** of the process.
- Solution: **add flag[2]** shared variable.
- If **flag[i] == true** → process i is **ready** to enter the critical section.
- Initial state:
flag[0] = flag[1] = false, turn = 0 or 1

Two-process solution

Process P_0

```
do {
```

```
    flag[0] = TRUE  
    turn = 1  
    while (flag[1] && turn==1){ }
```

critical section

```
    flag[0] = FALSE
```

remainder section

```
} while (TRUE)
```

Process P_1

```
do {
```

```
    flag[1] = TRUE  
    turn = 0  
    while (flag[0] && turn==0){ }
```

critical section

```
    flag[1] = FALSE
```

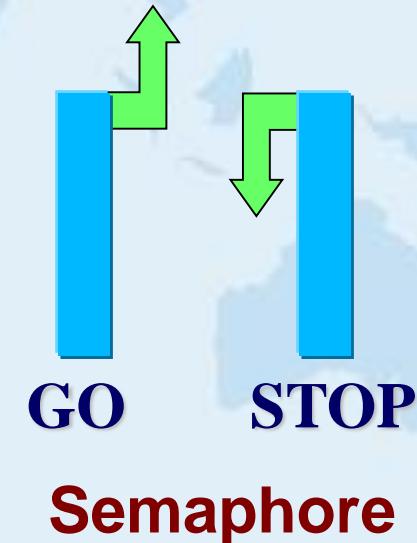
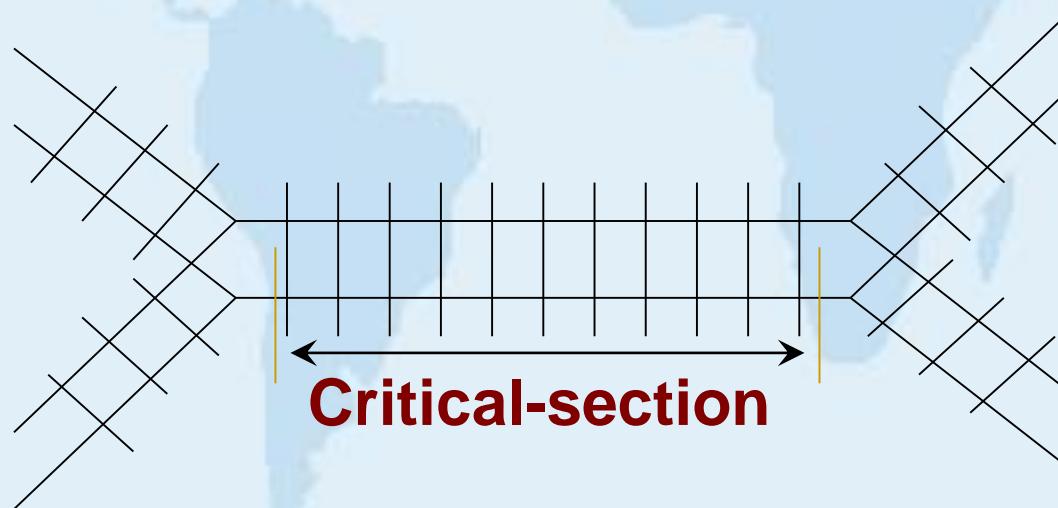
remainder section

```
} while (TRUE)
```

Process P_1 enters **critical section** when $\text{flag}[0] = \text{FALSE}$ or $\text{turn} = 1$.

Semaphores

- Previous solution is difficult to **generalize** to bigger problems, i.e. many processes.
- Another solution to the critical section problem is a **semaphore**.



Semaphore Operations

- A semaphore **S** is an integer variable, that, apart from initialization, is accessed **only** through two standard atomic operations: **wait** and **signal**.
- The classic definition of **wait** and **signal**:

```
wait(S) {  
    while S ≤ 0  
        do no-op  
    S--;  
}
```

```
signal(S) {  
    S++;  
}
```

Semaphore Implementation

- When one process modifies the semaphore value, no other process can simultaneously modify the same semaphore value.
- In addition, in case of **wait(S)**, the testing of integer value **S** ($S \leq 0$), and its possible modification (**S--**), must also be executed without interruption.

Semaphore: Usage

- A solution to the n -process critical-section problem by semaphore. The n processes share a semaphore, **mutex**, initialized to 1. Each process P_i is organized as follows:

```
do {
```

```
    wait (mutex)
```

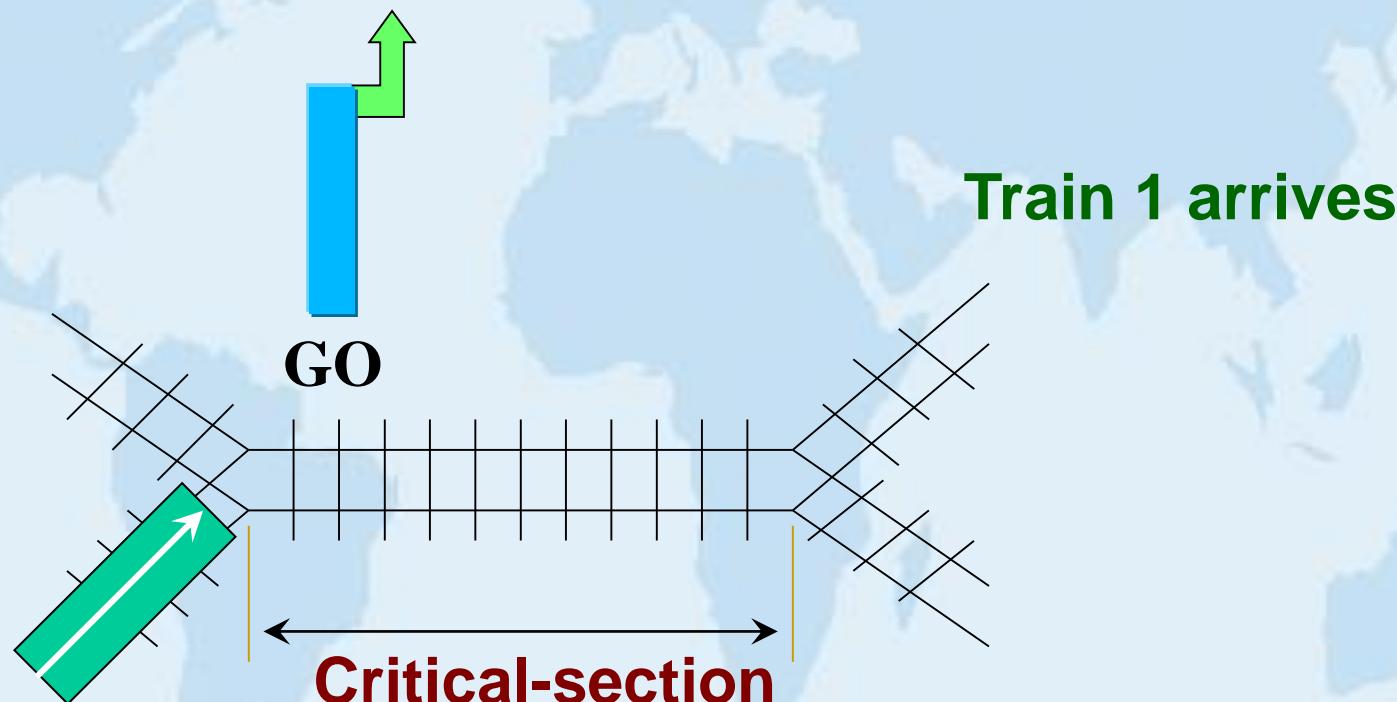
```
    // critical section
```

```
    signal (mutex)
```

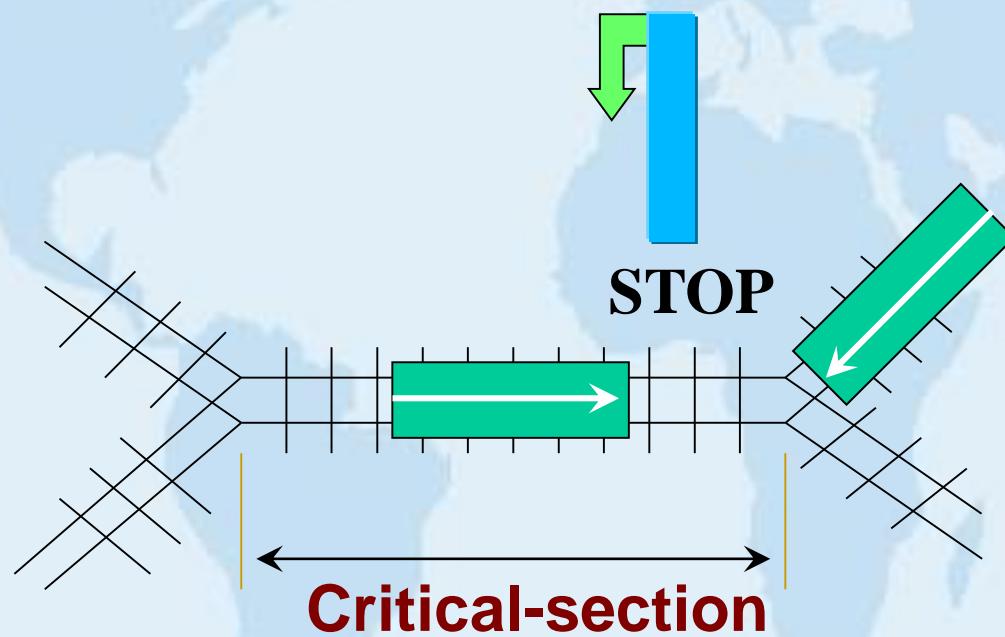
```
    // remainder section
```

```
} while (TRUE)
```

Semaphore Representation

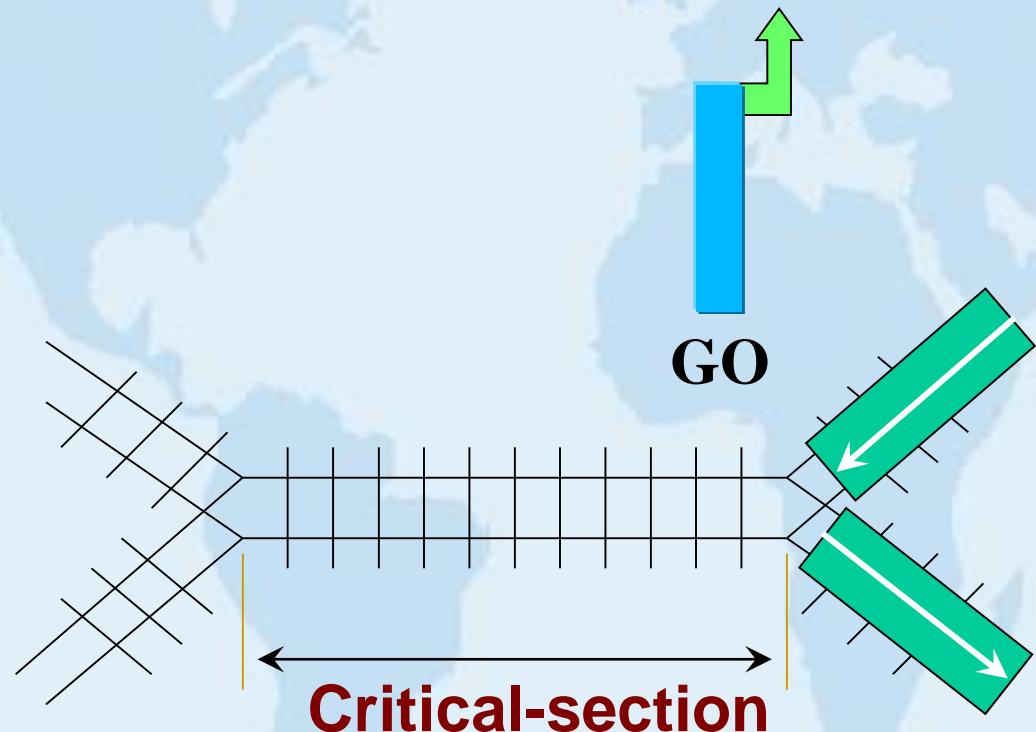


Semaphore Representation



Train 1 is in the
critical-section
and
train 2 arrives

Semaphore Representation



Train 1 lives the
critical-section and
train 2 can enter

Busy Waiting Problem

- The main disadvantage of the semaphore is the **busy waiting**: While a process is in its critical-section, any other process that tries to enter its critical-section must loop continuously in the wait statement, which wastes CPU time.
- This type of semaphore is called a **spinlock** (because process “spins” while waiting for lock).
- Waiting requires **no** context switch, so spinlocks (short time) are often used in multiprocessor systems.

Solution to busy waiting

- One solution to **busy waiting** is to associate each semaphore a **waiting queue**. When a process is in its critical-section, the other processes trying to enter their critical-sections are put in the waiting queue (process state changes from **running** to **waiting**).
- The control of CPU is transferred to the **CPU scheduler**. The FIFO or other strategies can be used for the queue associated to a semaphore.

Waiting state

- When some process executes a **signal** operation, a blocked process in the queue is restarted by **wakeup** operation (state changes from **waiting** to **ready**)
- Only one process is allowed executing **wait** and **signal** on the same semaphore at the same time. **Busy waiting** is useful for **wait** and **signal** in multiprocessor system.

Semaphore Type

- **Block** and **wakeup** involve context switch.

If context switch takes more time than spinlock, then **spinlock** is more efficient.

- Let's see a solution without **spinlock**.

To implement semaphores under this definition, we define a semaphore as:

```
typedef struct{
    int value;
    list Q;
} semaphore;
```

Semaphores operations

```
wait (semaphore S){  
    S.value--;  
    if S.value < 0 {  
        add this process to S.Q;  
        block (); }  
}  
  
signal (semaphore S):  
    S.value++;  
    if S.value ≤ 0{  
        remove process P from S.Q;  
        wakeup (P) ; }  
}
```

Semaphore Requirements

- Semaphores implementation must satisfy the following requirements:
 - ① **Mutual exclusion**
 - ② **Progress**
 - ③ **No deadlock**

Semaphore - Deadlock

- An example of deadlock with two semaphores S and Q.

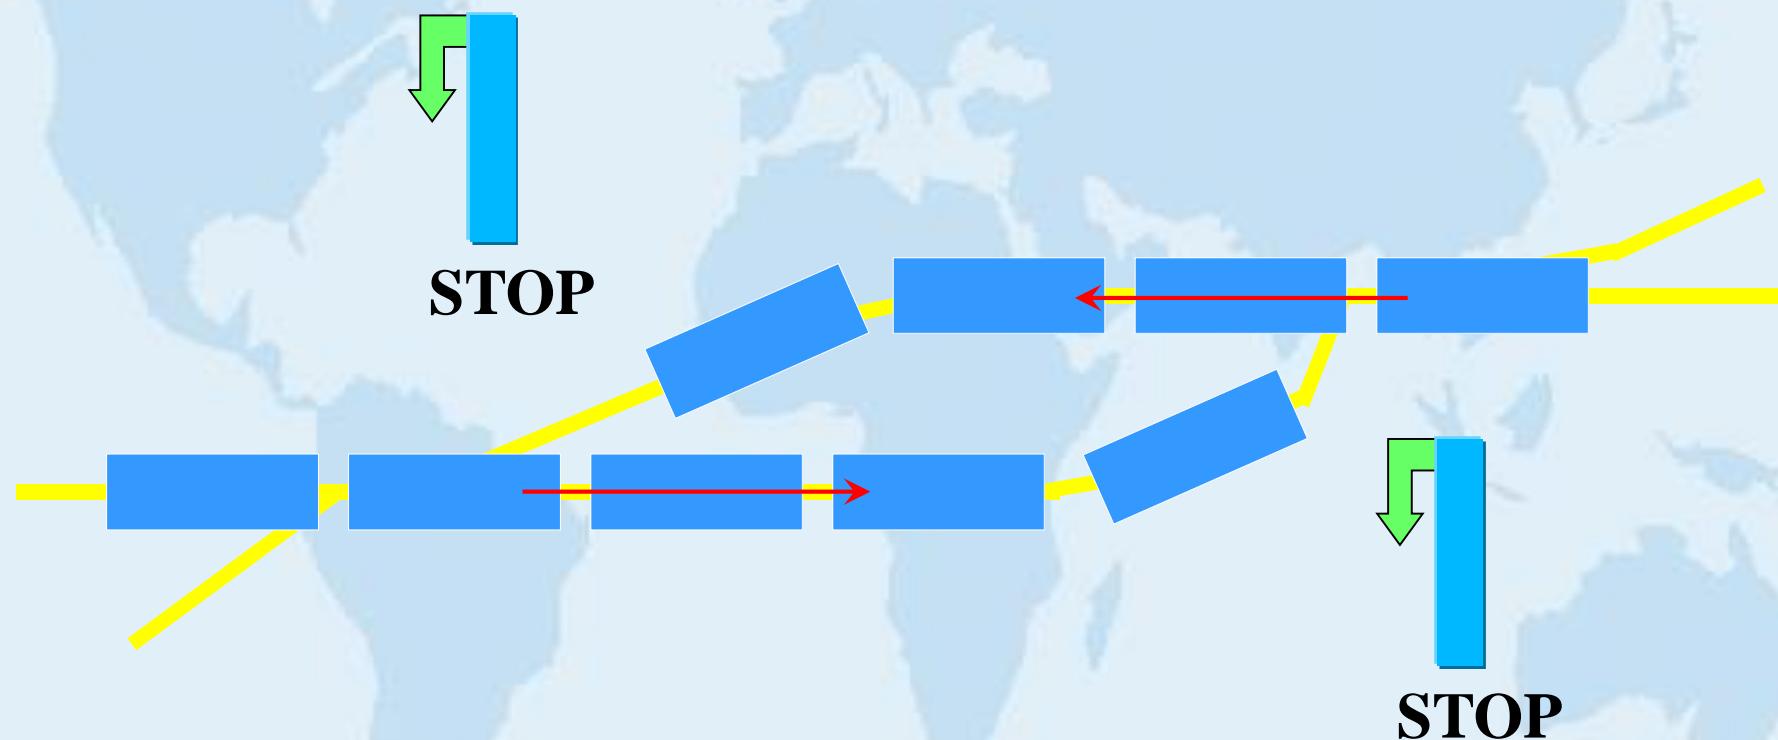
Process P0

```
Wait(S) ;  
Wait(Q) ;  
Critical section  
Signal(S) ;  
Signal(Q) ;
```

Process P1

```
Wait(Q) ;  
Wait(S) ;  
Critical section  
Signal(Q) ;  
Signal(S) ;
```

Semaphore - Deadlock



Synchronization Summary

- **Cooperating** process interact by sharing data and resources.
- To protect the shared access to data and resources, process must have **mutual exclusion** to critical-section code and must be **synchronized** (have an orderly execution).

Classical Problems

- **Classical problems** are simple examples that represent very complex, real problems.
- A classical problem shows the important features of the real problem and real solution.
- These problem are found in operating systems and applications.

Classical Problems

- ★ **The Bounded Buffer Problem**
(Producer-Consumer Problem)
- ★ **The Readers Writers Problem**
- ★ **The Dining Philosophers Problem**

The Bounded Buffer Problem

- Process P_0 is a **producer** of data.
- Process P_1 is a **consumer** of data.
- P_0 puts data into a **buffer** of bounded size n .
- P_1 gets the data from the buffer.
- If the buffer is **full**, P_0 waits.
- If the buffer is **empty**, P_1 waits.

The Bounded Buffer Problem Semaphores

- Three semaphores are used:
 - **mutex** (initialized to 0) provides mutual exclusion for access to the buffer.
 - **empty** (initialized to n) counts the empty locations of the buffer.
 - **full** (initialized to 0) counts the full locations of the buffer.
- **empty** and **full** are called **counting semaphores**.

The Bounded Buffer Problem Solution

Producer P_0

do {

```
produce data;  
wait(empty);  
wait(mutex);  
add data to  
    the buffer;  
signal(mutex);  
signal(full);  
}while (TRUE);
```

Consumer P_1

do {

```
wait(full);  
wait(mutex);  
remove data from  
    the buffer;  
signal(mutex);  
signal(empty);  
consume data;  
}while (TRUE);
```

The Readers Writers Problem

- Processes P_0, P_1, \dots, P_n **read** data from a file.
- Processes Q_0, Q_1, \dots, Q_m **write** data to the file.
- If a process reads data, other process **can** read data (**many readers OK**).
- If a process writes data, **no other** process can read or write data (**one writer OK**, read and write not OK).

The Readers Writers Problem Semaphores

- The **wrt** semaphore is for the mutual exclusion for access to the shared files.
- The **mutex** is for mutual exclusion for the count update.
- The **r_count** keeps the number of processes that are reading the file.
- **Starvation** of writers – when **r_count** does not reach 0!

The Readers Writers Problem Solution

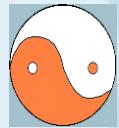
Writer Q

```
do {  
    wait(wrt) ;  
    writing;  
    signal(wrt) ;  
}while (TRUE) ;
```

Reader P

```
do {  
    wait(mutex) ;  
    r_count++ ;  
    if r_count=1; wait(wrt) ;  
    signal(mutex) ;  
    reading;  
    wait(mutex) ;  
    r_count-- ;  
    if r_count=0; signal(wrt) ;  
    signal(mutex) ;  
}while (TRUE) ;
```

The Dining Philosophers Problem



2 philosophers spend their lives thinking and eating. They share a table and 2 chopsticks (1 pair).

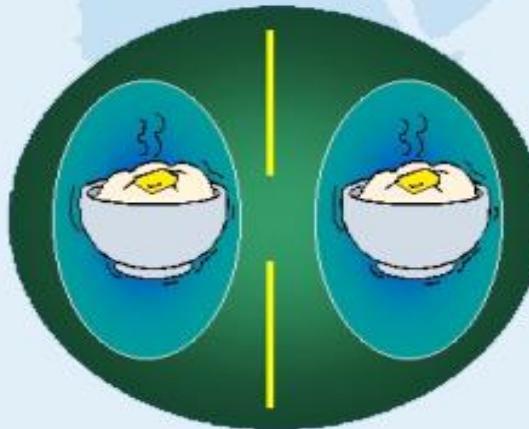
A philosopher gets 2 chopsticks to eat.

After eating, the philosopher puts down the 2 chopsticks.

Question: How can both philosophers eat?

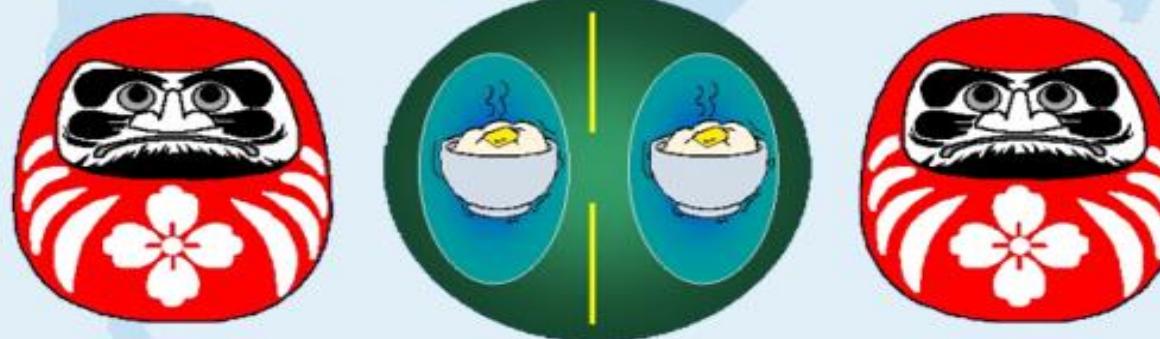
The Dining Philosophers Problem

- The situation of dining philosophers.



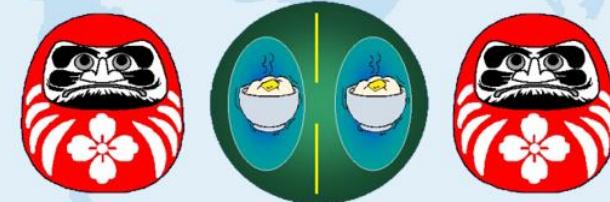
The Dining Philosophers Problem Semaphores

- Each chopstick is a semaphore.
- A philosopher **waits** for 2 chopsticks.
- A philosopher **signals** when finishing eating.



The Dining Philosophers Problem Solution

```
do {  
    wait(chopstick_1);  
    wait(chopsticks_2);  
    philosopher eats;  
    signal(chopstick_1);  
    signal(chopsticks_2);  
    philosopher thinks;  
} while (TRUE);
```



Summary

- For the cooperating process that share data, **mutual exclusion** must be provided.
- One solution is the ensure that a **critical-section** of code is in use by only one process at a time.
- The main disadvantage of these user coded solutions is **busy waiting**.
- **Semaphores** overcome this difficulty. They can be used to solve various synchronization problems.



That is all for today!