

# Operating systems

Lecture 5:  
**Deadlocks**

# System Model

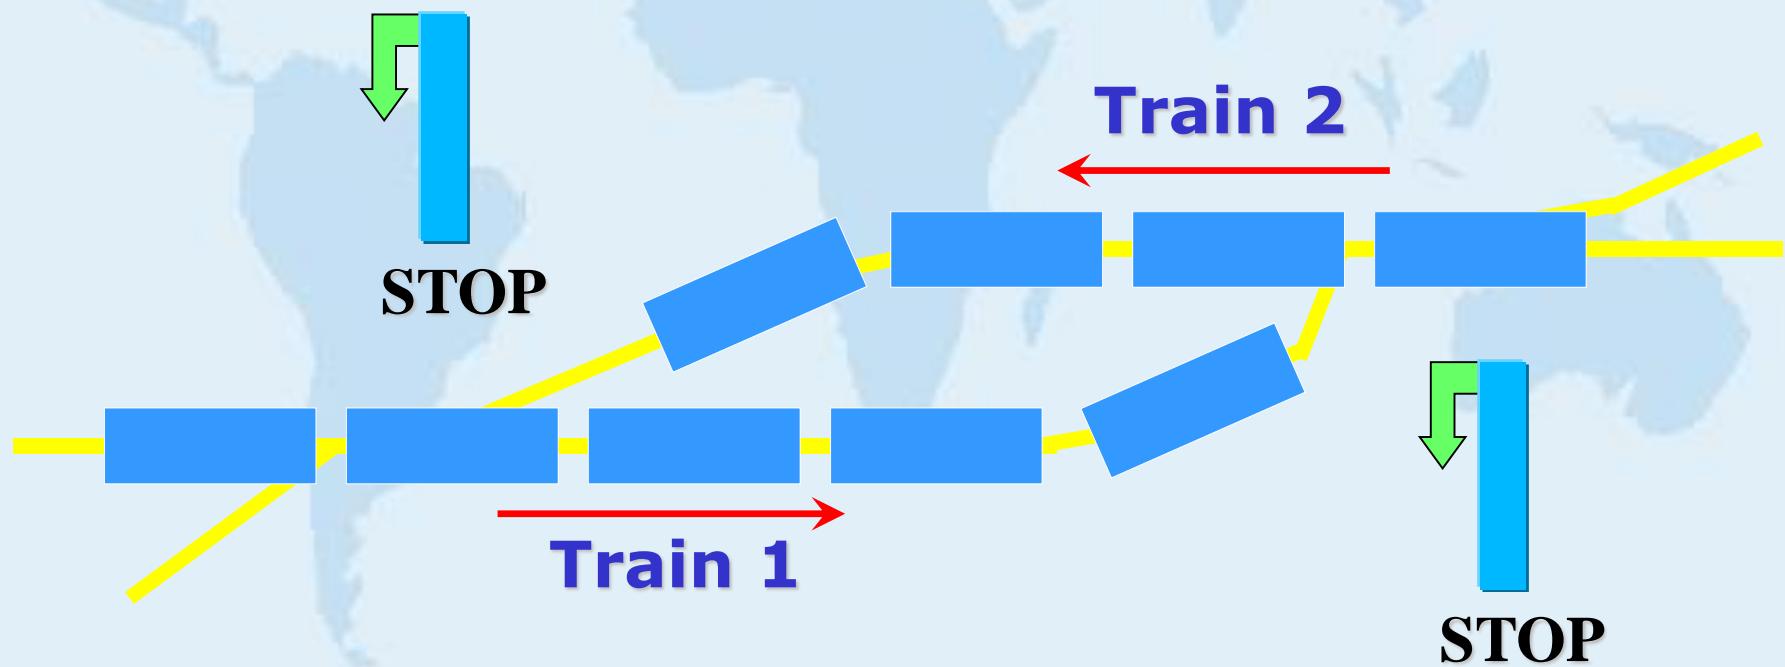
- Computer system consists of:
  - Number of **Processes** (programs in execution);
  - Finite number of **Resources** (CPU, memory space, files, I/O devices, etc.).
- Processes **compete** for resources.
  - If resource is not available – process waits.
  - **Deadlock** – when a process waits for resource held by another waiting process.

# System Model

- A process uses a resource in the following sequence:
  - **Request:** A process must request a resource before using it. If the resource is not available, the process enters waiting state.
  - **Use:** The process holds and operates on the resource.
  - **Release:** The process must release the resource after using it.
- The request and release are **System calls.**

# Deadlock example

A deadlock state occurs when **two or more** processes are waiting for a event that can be caused **only** by one of the waiting processes.



# Necessary Conditions

A **deadlock** situation can arise if the following four conditions hold **simultaneously** in a system:

## 1 **Mutual exclusion:**

At least one resource must be held in a non-sharable mode, i.e. only **one process** at a time can use the **resource**. If another process requests that resource, the requesting process must be delayed until the resource has been released.

# Necessary Conditions

## ② Hold and wait:

There must exist a process that is holding at least **one resource** and is **waiting** to acquire additional resources that are currently being held by other processes.

## ③ No preemption:

Resources **cannot** be preempted; that is, a resource can be released only **voluntary** by the process holding it, after that process has completed its task.

# Necessary Conditions

## ④ Circular wait:

There must exist a set  $\{P_0, P_1, \dots, P_n\}$  of waiting processes such that  $P_0$  is waiting for a resource that is held by  $P_1$ ,  $P_1$  is waiting for a resource that is held by  $P_2, \dots, P_{n-1}$  is waiting for a resource that is held by  $P_n$ , and  $P_n$  is waiting for a resource that is held by  $P_0$ .

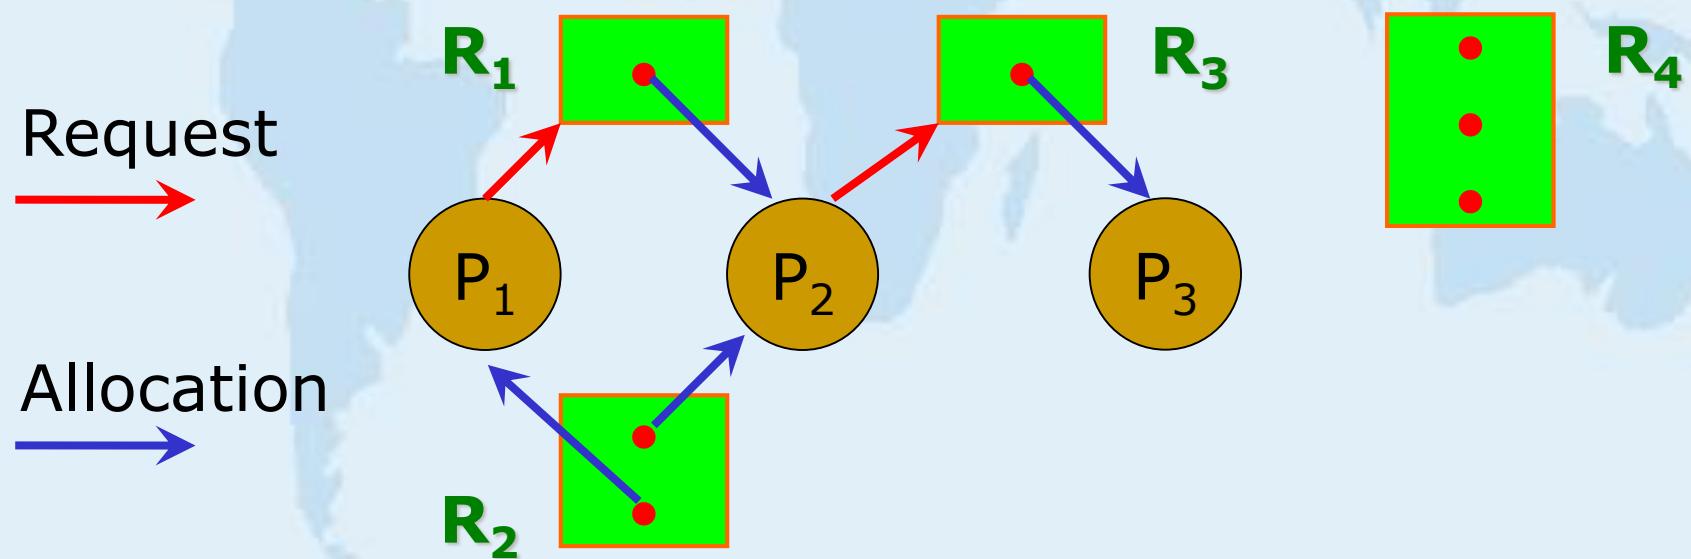
**For a deadlock to occur, all the four conditions must hold.**

# Resource-Allocation Graph

- Deadlocks can be described more precisely in terms of a directed graph, called a system **resource-allocation graph**.
- Nodes represent **processes**  $P = \{P_1, P_2, \dots, P_n\}$  and **resource types**  $R = \{R_1, R_2, \dots, R_m\}$ .
- Directed edge from process  $P_i$  to resource  $R_j$  is called **request edge**.
- Directed edge from resource  $R_k$  to process  $P_l$  is called **assignment edge**.

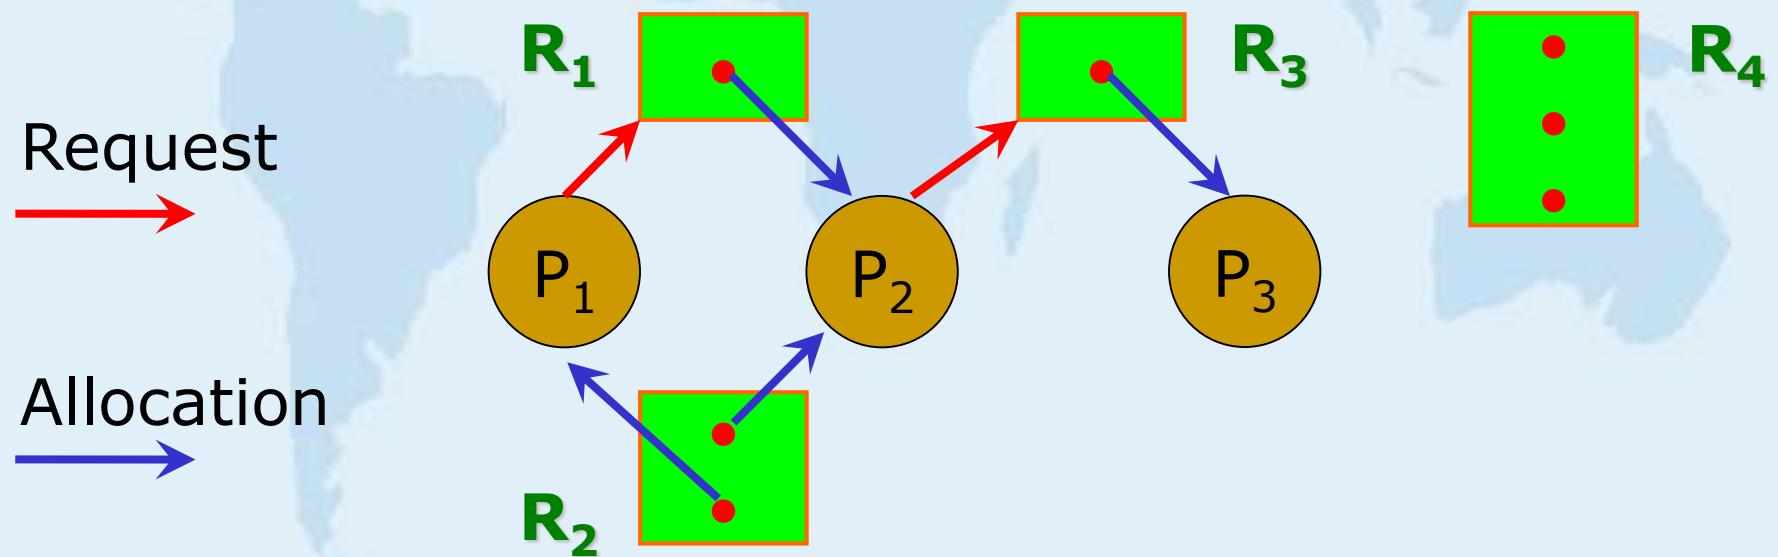
# Resource-Allocation Graph

- Processes are denoted as **circles** and resources are denoted as **squares**.
- Each **instance** of a resource is denoted as a dot in the square.



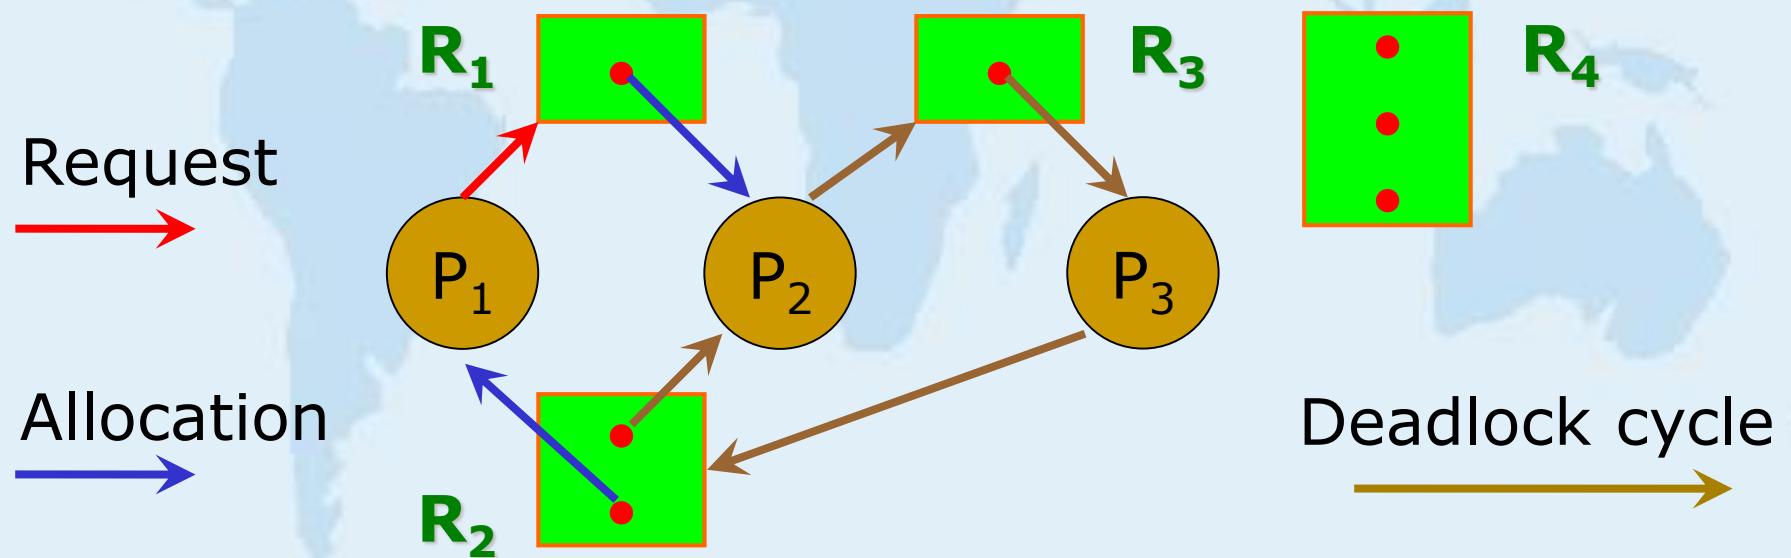
# Resource-Allocation Graph

- $P = \{P_1, P_2, P_3\}$        $R = \{R_1, R_2, R_3, R_4\}$
- $E = \{P_1 \rightarrow R_1, P_2 \rightarrow R_3, R_1 \rightarrow P_2, R_2 \rightarrow P_2, R_2 \rightarrow P_1, R_3 \rightarrow P_3\}$
- $R_1$  has 1 instance,  $R_2$  has 2,  $R_3$  has 1, and  $R_4$  has 3.
- Process  $P_1$  and  $P_2$  are in **waiting** state.



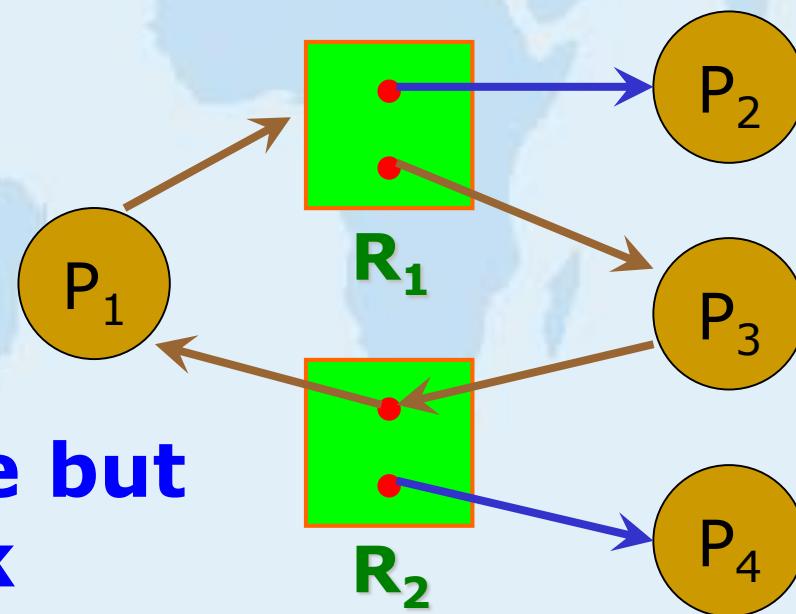
# Resource-Allocation Graph

- If the graph has no cycle, no process is deadlocked. If the graph has a **cycle**, deadlock may exist. If each resource type involved in a cycle has only **one instance**, the cycle implies a deadlock:



# Resource-Allocation Graph

- If each resource type involved in a cycle has **more** than one instance, the deadlock may or may not exist:



**With a cycle but  
no deadlock**

# Deadlock Handling Methods

Principally, there are three methods for dealing with the deadlock problem:

- ① We can use a **protocol** to ensure that the system will never enter a deadlock state.
- ② We can allow the system to **enter** a deadlock state and then **recover**.
- ③ We can **ignore** the problem all together, and pretend that deadlocks never occur in the system. (UNIX and many other OS.)

# Deadlock Handling Methods

- To ensure deadlocks do not occur, deadlock **prevention** or deadlock **avoidance** schemes can be used.
1. **Deadlock prevention** ensures that at least one of the four necessary conditions for deadlock **does not hold**.
  2. **Deadlock avoidance** requires information on resources that a process will use to be known **in advance** .

# Deadlock Prevention

It is necessary to **deny** one of the following:

- ① **Mutual exclusion:** In general cannot be denied because some resources are intrinsically non-sharable (e.g. printer).
- ② **Hold and wait:** Whenever process requests a resource, it should not hold any other resources.
  - Process requests and gets **all** the resources before it begins execution.
  - Leads to **lower** resource utilization and starvation.
  - Process requests resources **only** when it has none.
  - Waiting time **increases, low** process turnaround time.

# Deadlock Prevention

- ③ **No preemption:** Allow preemption of all resources held by processes which are waiting for other resources.
- ④ **Circular wait:** Define a total ordering of all the resources and require that each process requests resources in an increasing order of enumeration.

Assume  $T(R_i)$  is the order of  $R_i$ . After a process requests  $R_i$ , it can request another  $R_j$  only if  $T(R_j) > T(R_i)$ .

# Deadlock Prevention

**Disadvantages** of the Deadlock Prevention:

- **Restraints** how requests can be made.
- **Low** device utilization.
- **Reduced** system throughput.

# Deadlock Avoidance

- Methods for deadlock avoidance require **additional** information about how resources are going to be requested.
- The simplest model requires that each process declare the **maximum number** of resources that it may need.
- A **resource-allocation state** is defined by the number of **available**, **allocated**, and **maximum requested resources**.

# Deadlock Avoidance

- System is in **safe state** if there exists **safe sequence** of processes.
- A sequence of processes  $\langle P_1, P_2, \dots, P_n \rangle$  is a **safe sequence** for the current resource-allocation state if, for each  $P_i$ , the resources that  $P_i$  can still request can be satisfied by currently available resources plus the resources held by all the  $P_j$ , with  $j < i$ .
- If no safe sequence exists, system is in **unsafe state**.

# Deadlock Avoidance

- If resources that  $P_i$  needs are not immediately available, then  $P_i$  can **wait** until all  $P_j$  have finished.
- When all  $P_j$  finish and release their resources,  $P_i$  can obtain all of its needed resources, complete task, release them, and terminate.
- When  $P_i$  terminates,  $P_{i+1}$  can obtain its needed resources, and so on.

# Safe State - Example

(One resource class only)

- Total resources: 12
- Unallocated: 2

process	currently holds	max needed
A	4	6
B	4	11
C	2	7

- Safe sequence: A,C,B

# Unsafe State - Ex.

(One resource class only)

- Total resources: 12
- Unallocated: 2

process	currently holds	max needed
A	4	6
B	4	11
C	2	9

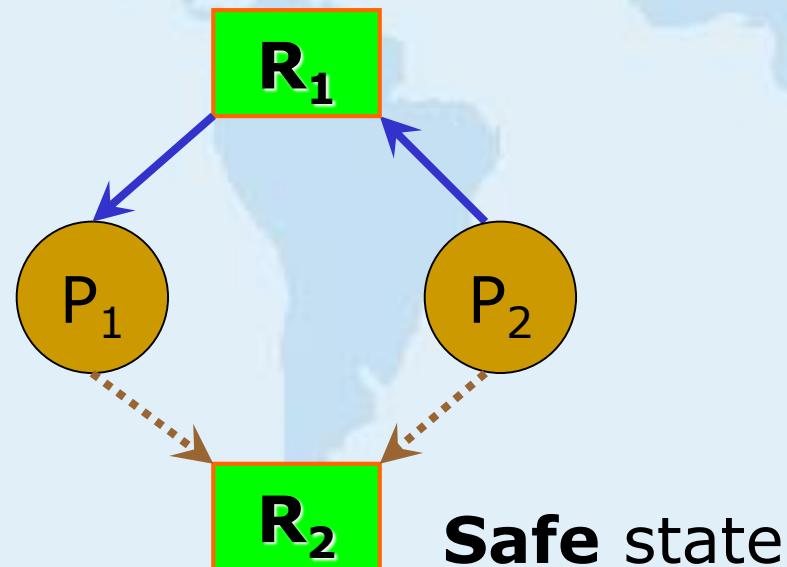
- **No** safe sequence exists!

# Resource-Allocation Graph

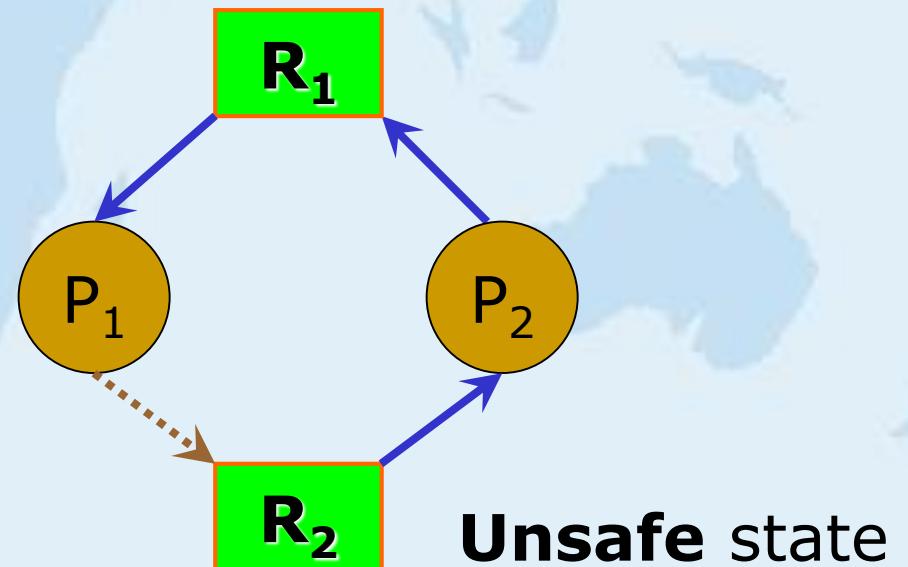
- Introduce new type of edge – **Claim edge**.
- Claim edge  $P_i \rightarrow R_j$  indicates that process  $P_i$  may request resource  $R_j$  in the future.
- Claim edge  $P_i \rightarrow R_j$  is represented by a dashed line.
- A claim edge  $P_i \rightarrow R_j$  can be added to the graph only if **all** the edges associated to  $P_i$  are claim edges.
- When  $P_i$  requests  $R_j$ , the **claim edge** is converted to a **request edge**.

# Resource-Allocation Graph

- Deadlock can be avoided if converting a request edge  $P_i \rightarrow R_j$  to an assignment edge  $R_j \rightarrow P_i$  **does not** produce a cycle in the graph.



**Safe state**



**Unsafe state**

# Deadlock Avoidance

## - Banker's Algorithm

- The resource-allocation graph algorithm is not applicable to a resource-allocation system with multiple instances for each resource type.
- For this case, we will use the **banker's algorithm**.
- This algorithm could be used in a banking system to ensure that the bank never allocates its available cash such that it can no longer satisfy the needs of all its customers.

# Deadlock Avoidance

## - Banker's Algorithm

- When a new process enters the system, it must **declare** the maximum number of each resource type that it may need. This number may not exceed the total number of resources.
- When a user requests a set of resources, the system must **determine** whether the allocation of these resources will leave the system in a safe state.
- If it will, the resources are allocated; otherwise, the process must **wait** until some other process releases enough resources.

# Deadlock Avoidance: Banker's Algorithm

- Let  $n$  be the number of processes in the system and  $m$  be the number of resource types. Define following:
- ✓ **Available**:  $\text{Avail}[1..m]$ .  $\text{Avail}[j] = k$  means  $k$  instances of  $R_j$  are available.
- ✓ **Max**:  $\text{Max}[1..n, 1..m]$ .  $\text{Max}[i, j] = k$  means  $P_i$  requests at most  $k$  instances of  $R_j$ .
- ✓ **Allocation**:  $\text{Alloc}[1..n, 1..m]$ .  $\text{Alloc}[i, j] = k$  means  $P_i$  holds  $k$  instances of  $R_j$ .
- ✓ **Need**:  $\text{Need}[1..n, 1..m]$ .  $\text{Need}[i, j] = k$  means  $P_i$  may need  $k$  additional instances of  $R_j$ .

# Deadlock Avoidance: Banker's Algorithm

- ✓ **Request:**  $\text{Req}[1..n, 1..m]$ .  $\text{Req}[i, j] = k$  means  $P_i$  wants  $k$  instances of resource type  $R_j$ .
- ✓ Note that  $\text{Need}[i, j] = \text{Max}[i, j] - \text{Alloc}[i, j]$ .
- ✓ In addition:
  - ✓ For  $X[1..n]$  and  $Y[1..n]$ ,  $X \leq Y$  if  $X[i] \leq Y[i], 1 \leq i \leq n$ .
  - ✓  $X < Y$  if  $X \leq Y$  and  $X \neq Y$ .
- ✓ The  $i$ -th row of  $\text{Max}$ ,  $\text{Alloc}$ ,  $\text{Need}$  is denoted as  $\text{Max}_i$ ,  $\text{Alloc}_i$ ,  $\text{Need}_i$ , respectively.

# Safe state checking algorithm

## 1. Initialization

$Work[1..m] = Avail [1..m]$ ;  $Finish[i] = \text{false}$ ,  $1 \leq i \leq n$ .

## 2. Find $i$ such that

$Finish[i] = \text{false}$  and  $\text{Need}_i \leq Work$

then go to step 3 else go to step 4;

## 3. Execute

$Work = Work + Allo_i$ ;  $Finish[i] = \text{true}$ ;

go to step 2;

## 4. Finish

If  $Finish[i] == \text{true}$  for all  $i$ , then the system is **safe**.

# Resource-Request Algorithm

When a request for resources is made by process  $P_i$ :

1. **If**  $Req_i \leq Need_i$  **then** go to step 2 **else** raise error  
(since the process has exceeded its maximum claim).
2. **If**  $Req_i \leq Avail$  **then** go to step 3 **else**  $P_i$  must wait  
(since the resources are not available).
3. **Modify** state as follows:  
 $Avail = Avail - Req_i$   
 $Alloc_i = Alloc_i + Req_i$   
 $Need_i = Need_i - Req_i$
4. **Check** if state is safe (call previous algorithm) and then allocate resources to  $P_i$

# Banker's Algorithm: An Example

- Five processes  $P_0, P_1, P_2, P_3, P_4$ . Three types of resources **A** (10 instances), **B** (5 instances), **C** (7 instances). At time  $T_0$  we have:

Avail	Alloc	Max	Need					
A	B	C	A	B	C	A	B	C
3	3	2						
$P_0$	0	1	7	5	3	7	4	3
$P_1$	2	0	3	2	2	1	2	2
$P_2$	3	0	9	0	2	6	0	0
$P_3$	2	1	2	2	2	0	1	1
$P_4$	0	0	4	3	3	4	3	1

**Safe state** sequence:  
(by the algorithm)  
 $P_1, P_3, P_4, P_2, P_0$

# Banker's Algorithm: An Example

- Assume  $P_1$  gives additional request  $Req_1 = (1,0,2)$ . We first check  $Req_1 \leq Avail$ :  $(1,0,2) \leq (3,3,2)$ . The system enters following state:

Avail	Alloc	Max	Need
A B C	A B C	A B C	A B C
2	3	0	
$P_0$	0 1 0	7 5 3	7 4 3
$P_1$	3 0 2	3 2 2	0 2 0
$P_2$	3 0 2	9 0 2	6 0 0
$P_3$	2 1 1	2 2 2	0 1 1
$P_4$	0 0 2	4 3 3	4 3 1

**Safe state** sequence:  
(by the algorithm)  
 $P_1, P_3, P_4, P_2, P_0$

# Banker's Algorithm: An Example

- By the safety algorithm,  $P_1, P_3, P_4, P_2, P_0$  is a safe sequence and  $\text{Req}_1$  can be granted.
- If  $P_4$  requests  $(3,3,0)$  at this state, the request **cannot** be granted since the resources are not available.
- If  $P_0$  requests  $(0,2,0)$ , even though the resources are available, the granting will make the system **unsafe**.

# Deadlock Detection

If a system **doesn't** have neither a deadlock-prevention nor a deadlock-avoidance, then a deadlock **may** occur. In this case, the system **must provide**:

- An algorithm that examines the state of the system to **determine** whether a deadlock has occurred.
- An algorithm to **recover** from a deadlock.

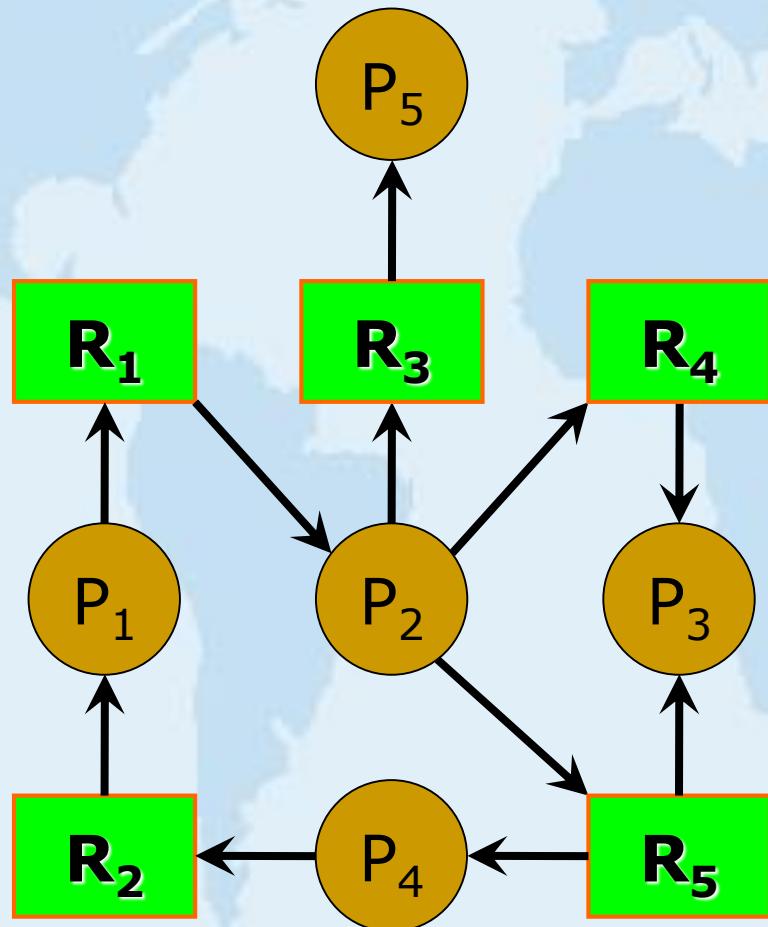
# Deadlock Detection: Single Instance Case

When **all** resources have only a **single** instance:

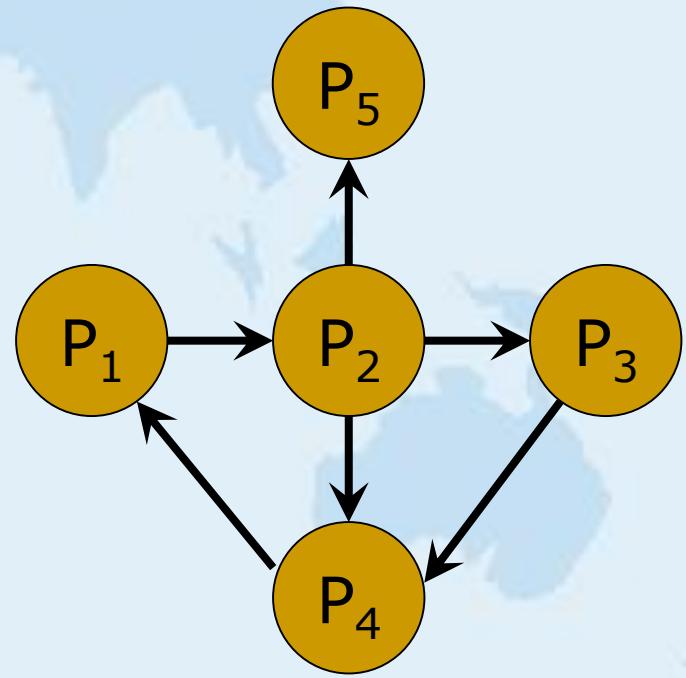
- Define a deadlock detection algorithm that uses a variant of the **resource-allocation** graph, called a **wait-for** graph.
- This graph is obtained from the resource-allocation graph by **removing** the resource nodes and collapsing the appropriate edges.

# Deadlock Detection: Single Instance Case

Resource-allocation graph



Corresponding wait-for graph



# Deadlock Detection: Single Instance Case

- A deadlock **exists** in the system **if and only if** the **wait-for** graph contains a cycle.
- To detect deadlocks, the system needs to maintain the **wait-for** graph and periodically invoke an algorithm that **searches** for a cycle in the graph.
- An algorithm requires  **$O(n^2)$**  operations, where  **$n$**  is the number of vertices in the graph.

# Deadlock Recovery: Process Termination

To recover from deadlock we can **terminate**:

- All deadlocked processes - Simple method, high expense.
- Abort one process at a time until the deadlock cycle eliminated - Large overhead, since, after each process is aborted, **deadlock-detection** algorithm must be invoked to determine whether any process are still deadlocked.

# Deadlock Recovery: Process Termination

To **choose** which process to abort, the following should be considered:

- The **priority** of the process.
- The **time** the process has executed and will be executed.
- How many and what type resources are **used**.
- Resources **needed** further.
- **Number** of processes needed to be aborted.
- Whether the process is **interactive** or **batch**.

# Deadlock Recovery: Resource Preemption

To eliminate deadlocks we can use **resource preemption**:

- Successively **preempt** some resources from processes and **give** these resources to other processes until the deadlock cycle is broken.
- If preemption is required to deal with deadlocks, then some **issues** need to be addressed.

# Deadlock Recovery: Resource Preemption

- 1. Selecting a victim** - Which resources and which processes are to be preempted?  
Decide the order of preemption.
- 2. Rollback** - If we preempt a resource from a process, we must roll back the process to some safe state, and restart it.
- 3. Starvation** - How to prevent starvation?  
That is, how we guarantee that resources will not always be preempted from the same process?

# Deadlock: Combined Approach

To handle deadlock situations **combinations** of the following approaches can be used:

- **Deadlock-prevention**,
- **Deadlock-avoidance**, and
- **Detection-recovery**.

# Deadlock: Summary

- A deadlock state occurs when two or more processes are waiting **indefinitely** for an event that can be **caused** only by one of the waiting processes.
- Three methods for handling deadlocks:
  - 1) To ensure the system **never** enters a deadlock state.
  - 2) Allow the system to **enter** the deadlock state and then **recover**.
  - 3) **Ignore** the problem.

# Deadlock: Summary

- **Deadlock-prevention** ensures one of the conditions never occur.
- **Deadlock-avoidance** requires a priori information about the requested resources of each process.
- **Detection and recovery** uses wait-for graphs for detection. Termination of processes and preemption of resources can be used in deadlock recovery.



**That is all for today!**