

# Operating systems

## **Lecture 7:**

## **Memory management**

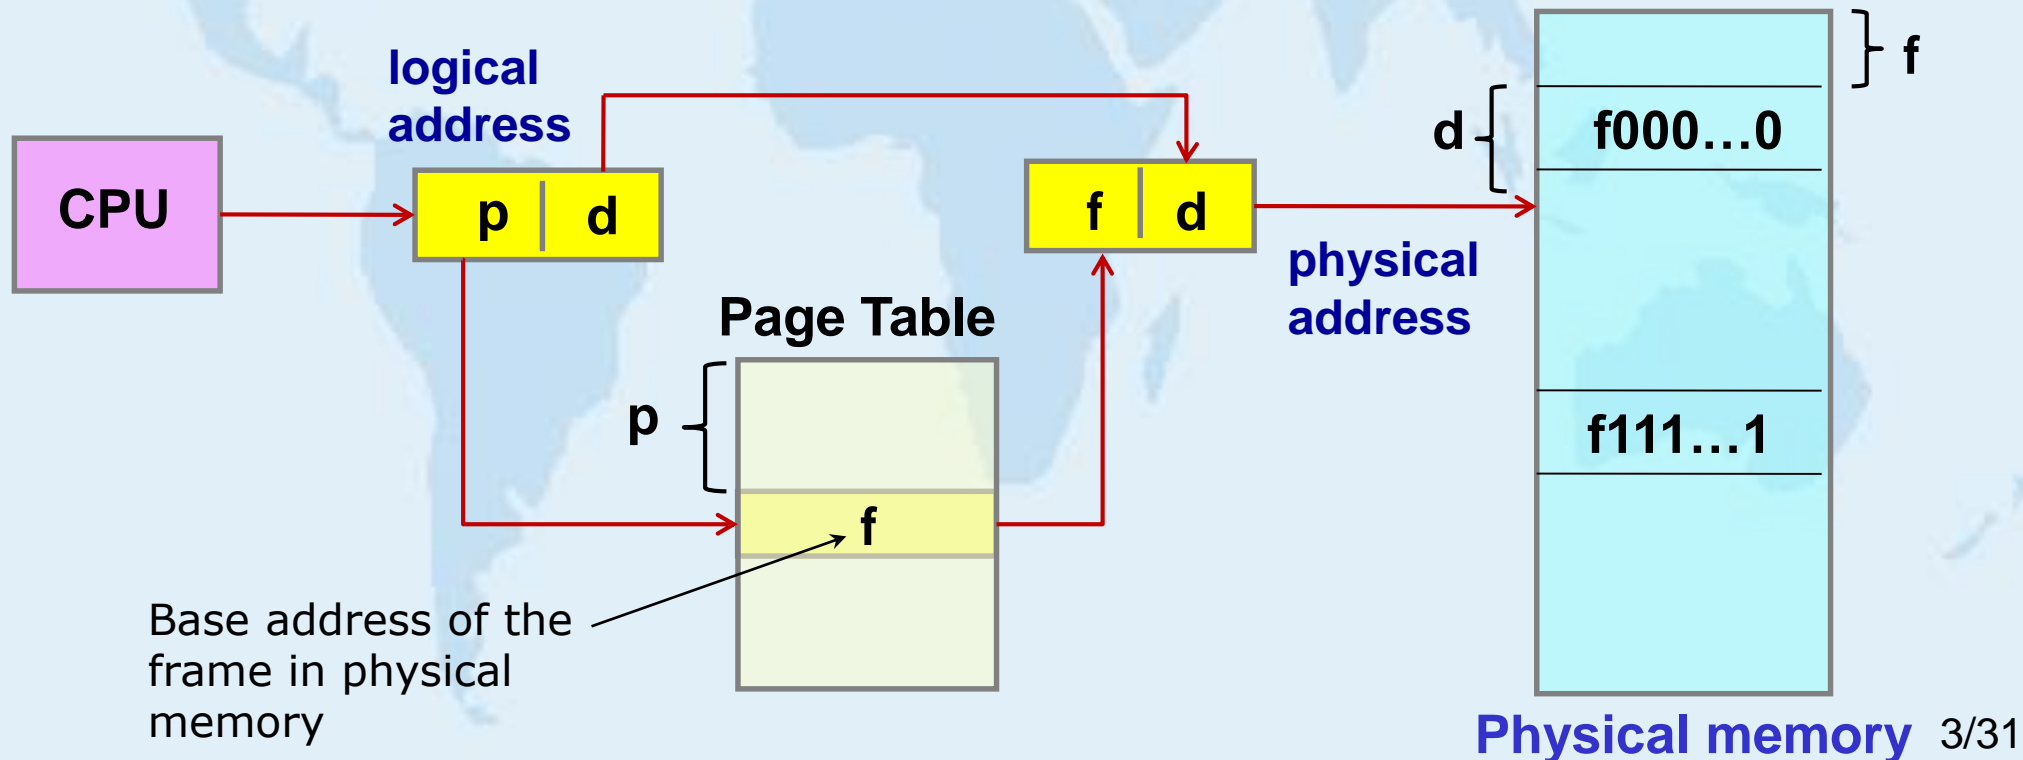
- Paging**
- Segmentation**

# Paging

- **Paging** – memory management scheme that allows physical address space to be non-contiguous.
- **Basic idea:**
  - Divide logical memory into fixed-sized blocks (size is power of 2, between 512 and 8192 bytes) called **pages**.
  - Divide physical memory into blocks of the same size called **page frames**.
  - Keep track of all free page frames.
  - To run a program of size **n** pages, need to find **n** free page frames and load the program.

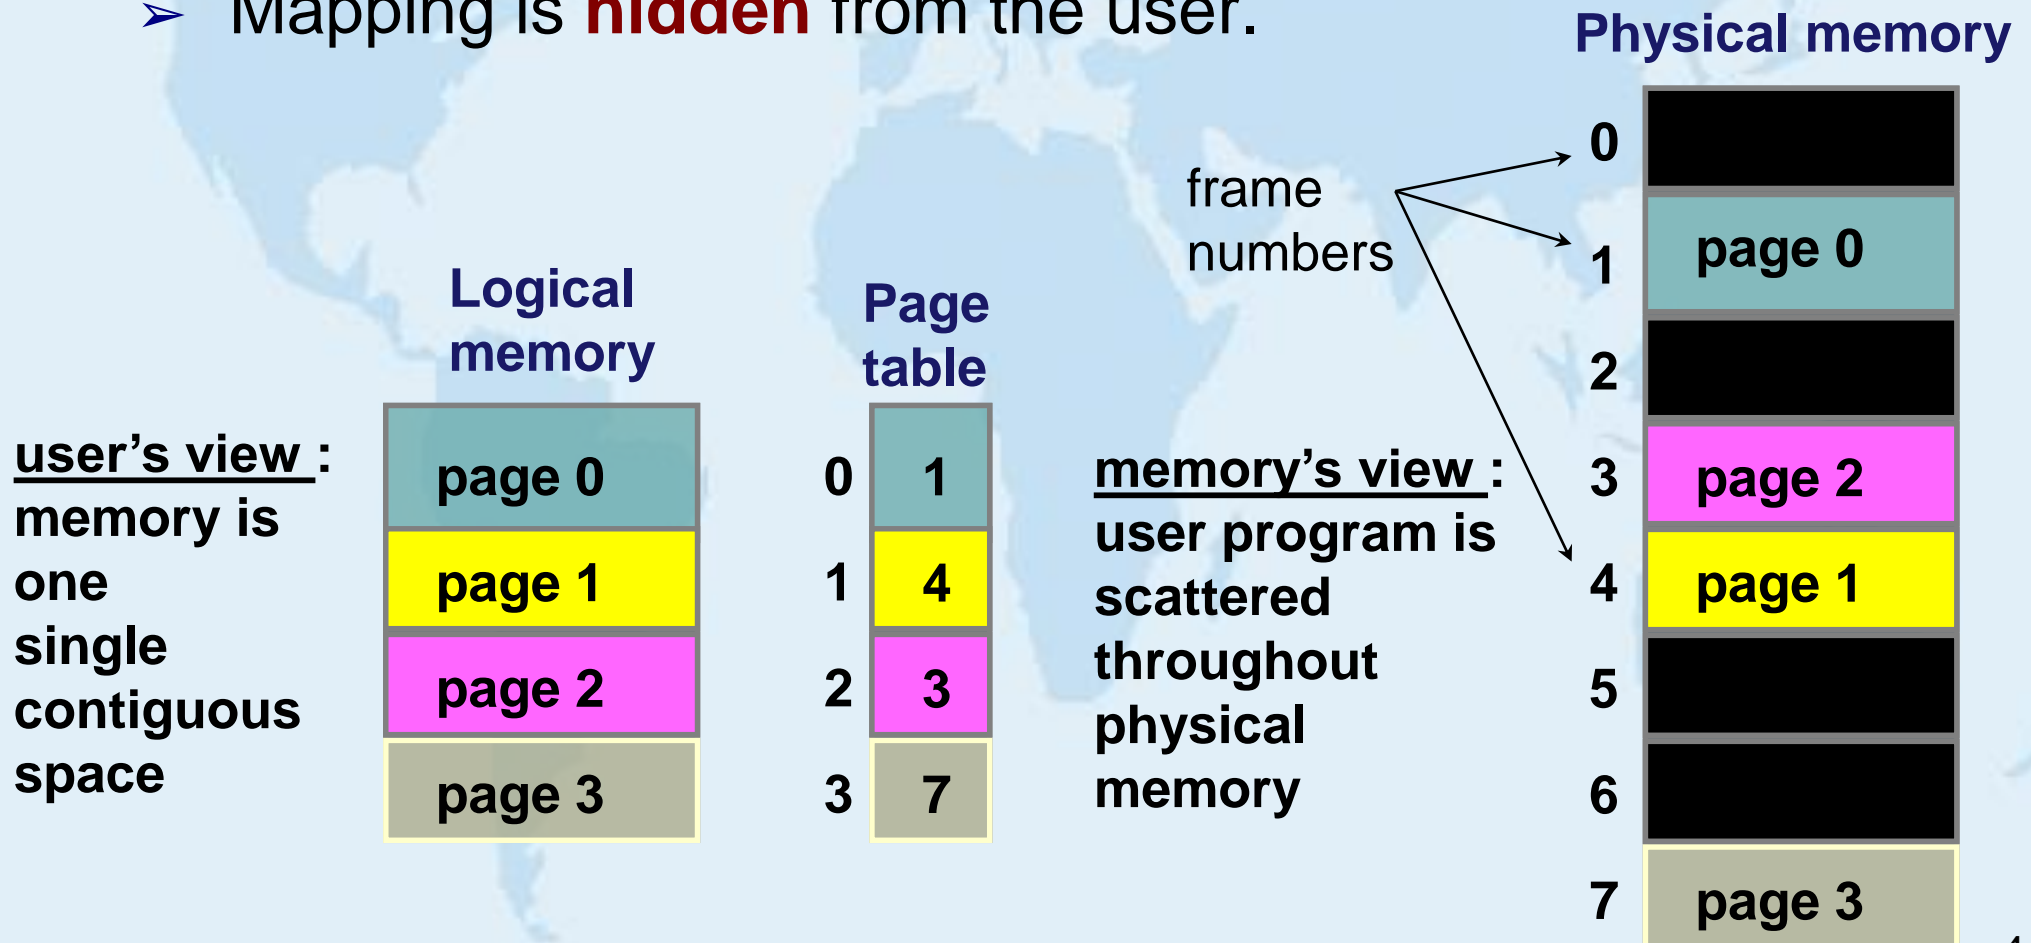
# Paging

- Addresses generated by the CPU are divided into:
  - Page number (p)** - used as an index into a **Page Table** which contains base address of each page in the physical memory.
  - Page offset (d)** - combined with base address to define the physical memory address that is sent to the memory unit.

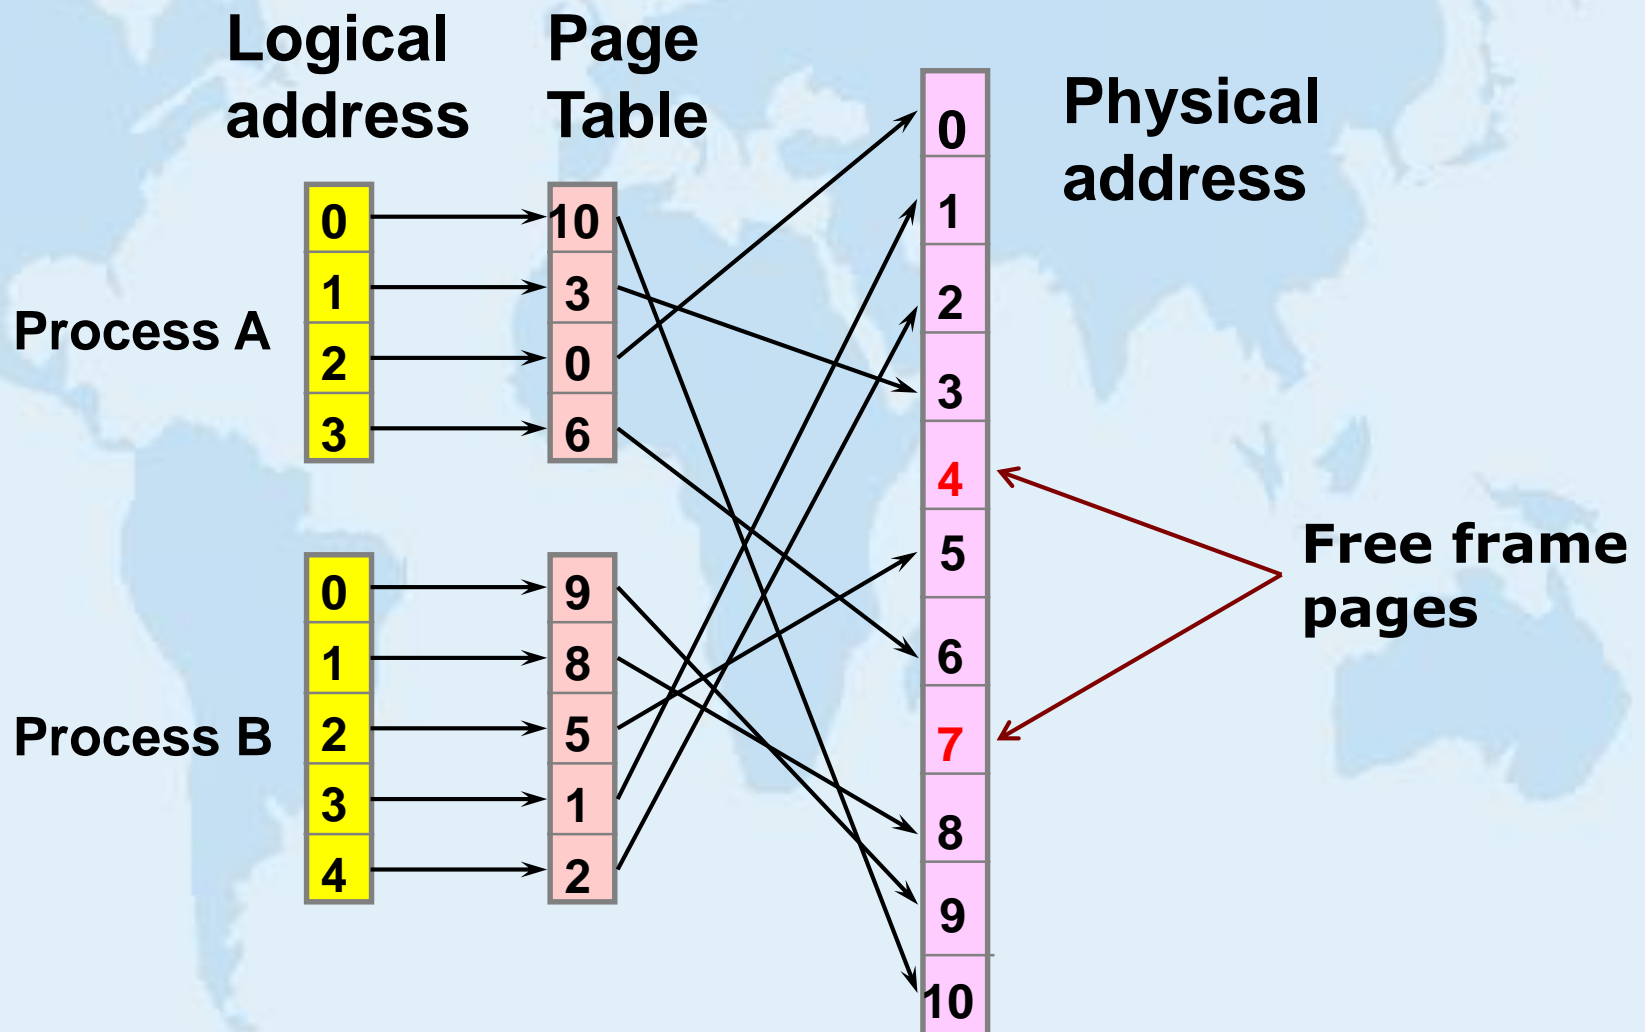


# Paging

- **Logical-to-physical** address mapping uses paging.
- Mapping is **hidden** from the user.




# Paging Example




# Page Size

**Page size = 1kB**




**Process  
with  
5.5 kB**



Internal  
fragmentation

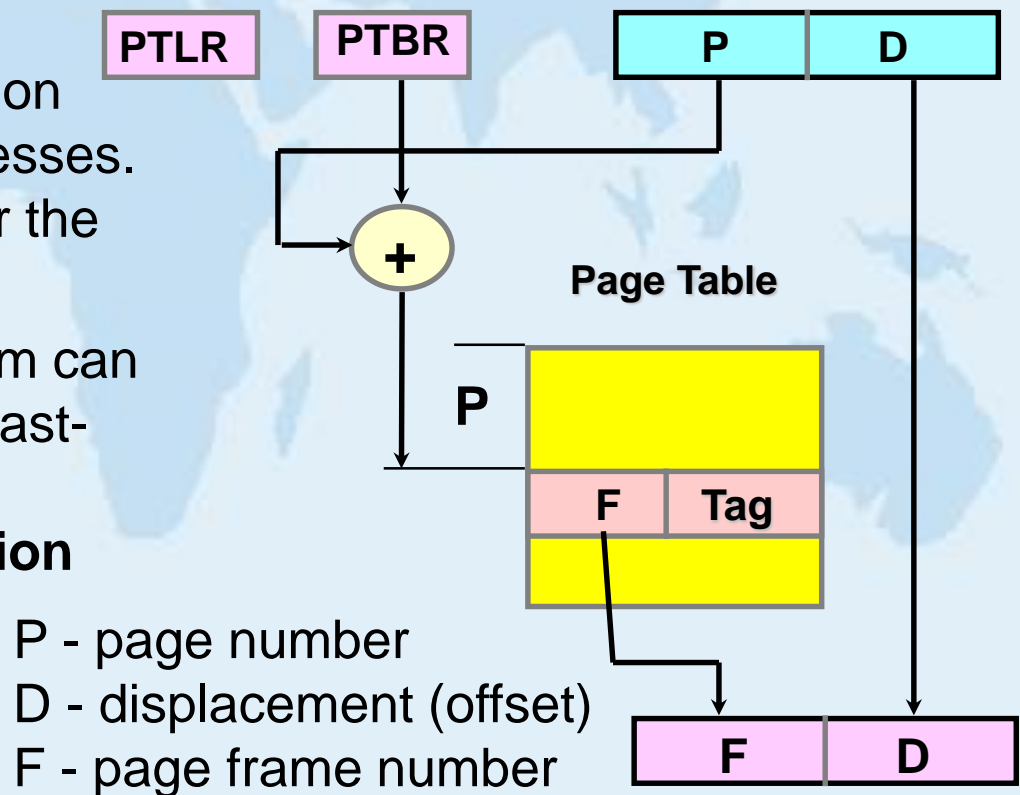
**Page size = 4kB**



# Page Table

- Page table is kept in **main** memory. One table per process.
- **Page-table base register (PTBR)** points to the page table.
- **Page-table length register (PTLR)** indicates size of page table.

- In this scheme every data/instruction access requires **two** memory accesses. One for the page table and one for the data/instruction.
- The two memory accesses problem can be solved by the use of a special fast-lookup hardware cache called **associative registers** or **translation look-aside buffers (TLB)**.

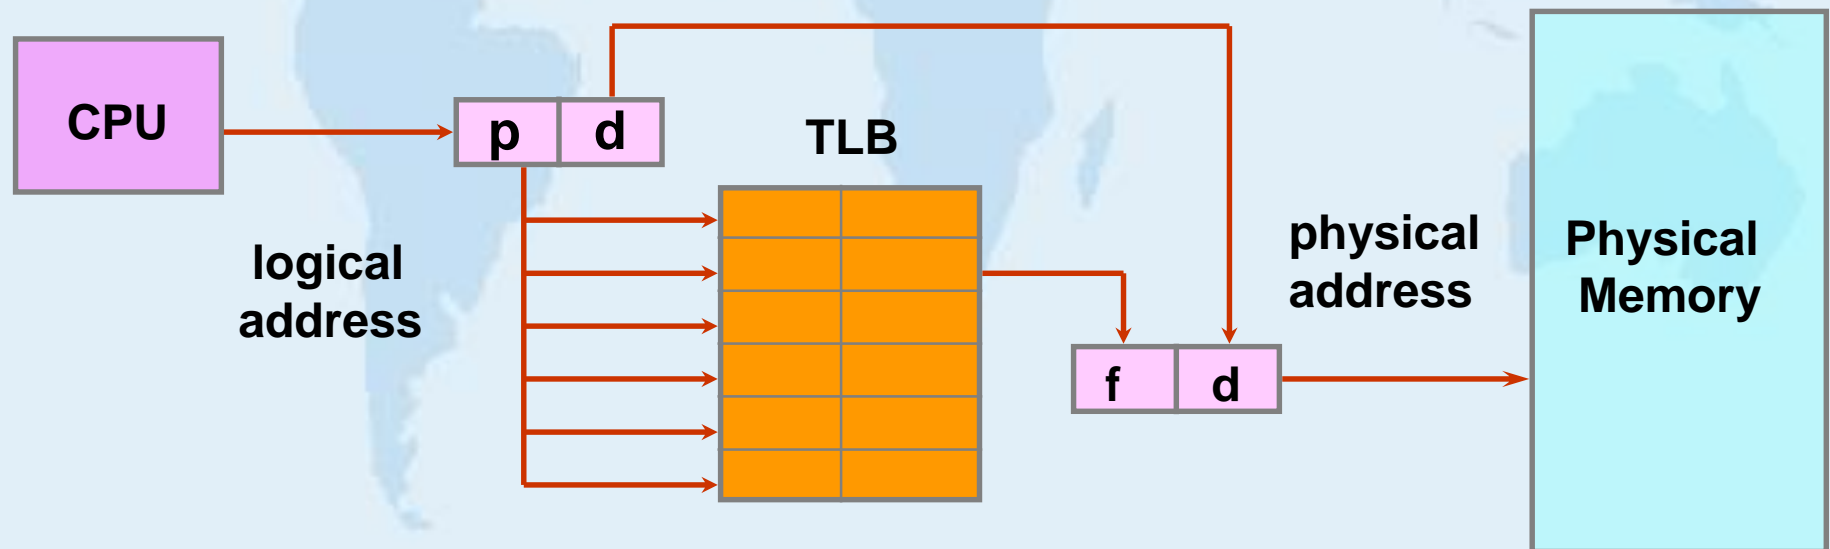




# Associative Memory

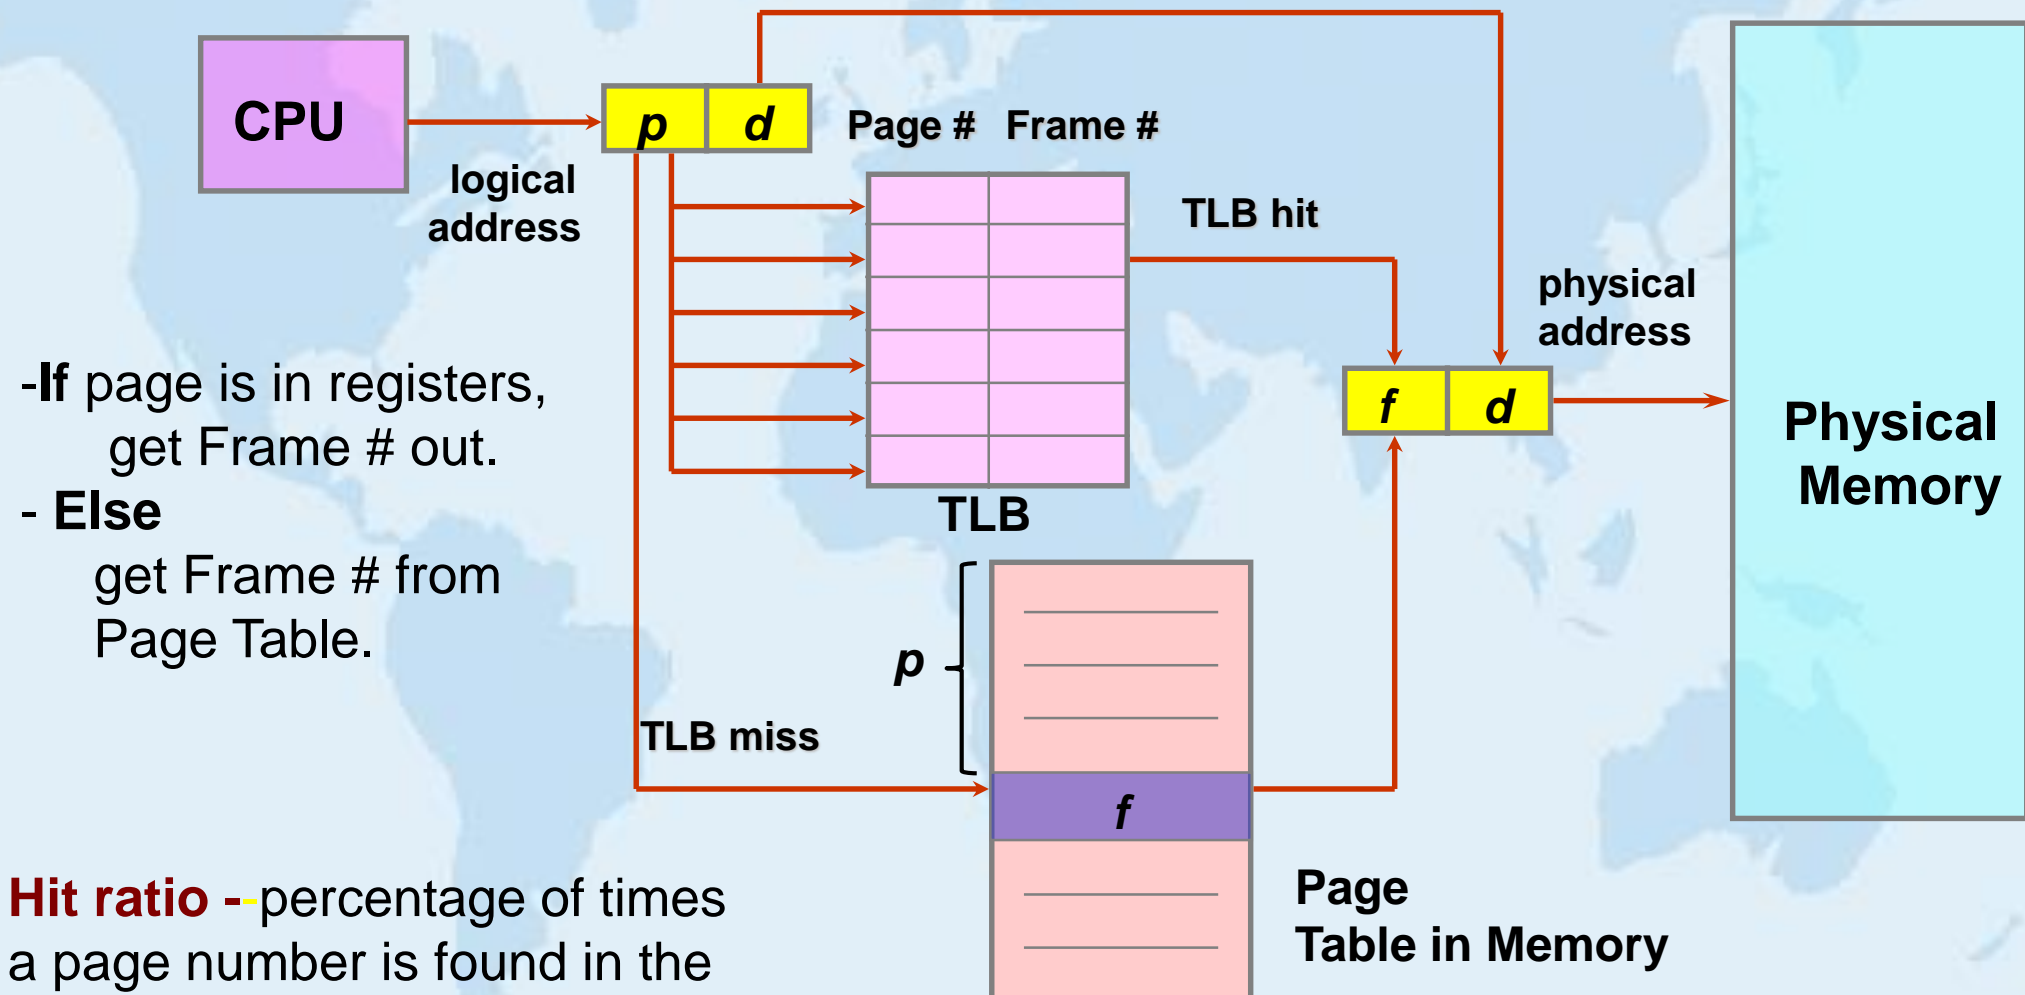
## TLB

- **Associative memory** - Memory organized in such a way that memory cells can be searched (in parallel) by contents, not the address or index.
- Implemented with a high-speed logic such as **registers**.
- Used to store and read a **small** amount of frequently used information (64 – 1024 bytes).





# TLB Paging Hardware



- If page is in registers, get Frame # out.
- Else get Frame # from Page Table.

**Hit ratio** - percentage of times a page number is found in the registers; this ratio is related to number of associative registers.

# Effective Access Time

## ➤ Example:

- TLB access time:  $T_B$  ( $\sim 20$  ns)
- Memory access time:  $T_M$  ( $\sim 100$  ns)
- Hit ratio:  $P_{HIT}$  (80%  $\sim$  98%)

## ➤ Effective Memory Access Time (**EAT**):

$$T_E = P_{HIT} (T_B + T_M) + P_{MISS} (T_B + T_M + T_M)$$

$$P_{HIT} + P_{MISS} = 1$$

$$P_{hit} = 80\% \quad \mathbf{EAT} = 0.80 \times 120 + 0.20 \times 220 = 140 \text{ ns (40\% slowdown)}$$

$$P_{hit} = 98\% \quad \mathbf{EAT} = 0.98 \times 120 + 0.02 \times 220 = 122 \text{ ns (22\% slowdown)}$$

# Memory Protection

- Protection **bits** are associated with each frame.
  - **Read-write** – process can read and write to this frame.
  - **Read-only** – process can only read from this frame
  - **Valid-invalid** -
    - “**valid**” indicates that the associated page is in the process’ logical address space, and is thus a legal page.
    - “**invalid**” indicates that the page is not in the process’ logical address space.

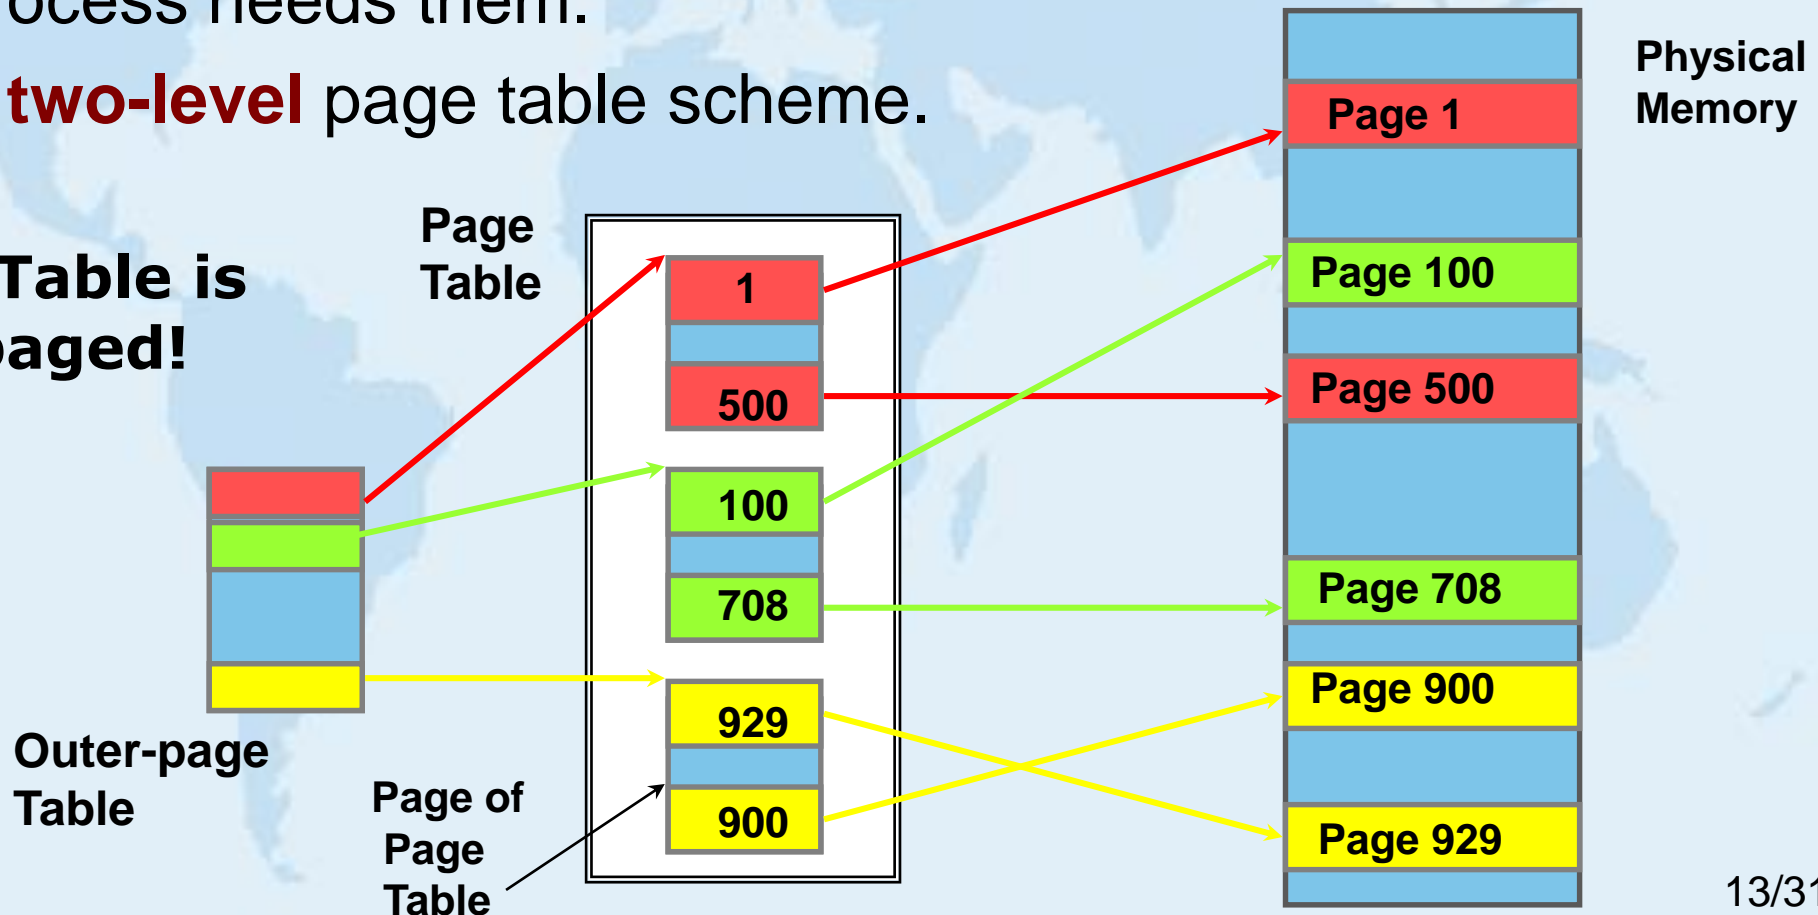
# Avoiding Table Fragmentation

- Each process has a Page Table which **takes** memory:
  - 32-bit logical address space 4GB ( $2^{32}$ ).
  - Page size is 4K bytes ( $2^{12}$ ).
  - Page Table may consist of up to 1 M entries ( $2^{32} / 2^{12}$ ).
    - Each entry is 4 bytes => each process needs up to 4 Mbytes of physical address space for the Page Table.
- **Problem** – Need 4M of **contiguous** main memory!
- Solutions:
  - **Multilevel Paging.**
  - **Hashed Page Table.**
  - **Inverted Page Table.**

# Multilevel Paging

- **Multilevel Paging** - partitioning the Page Table allows the operating system to leave partitions unused until a process needs them.
- A **two-level** page table scheme.

**Page Table is also paged!**



# Multilevel Paging

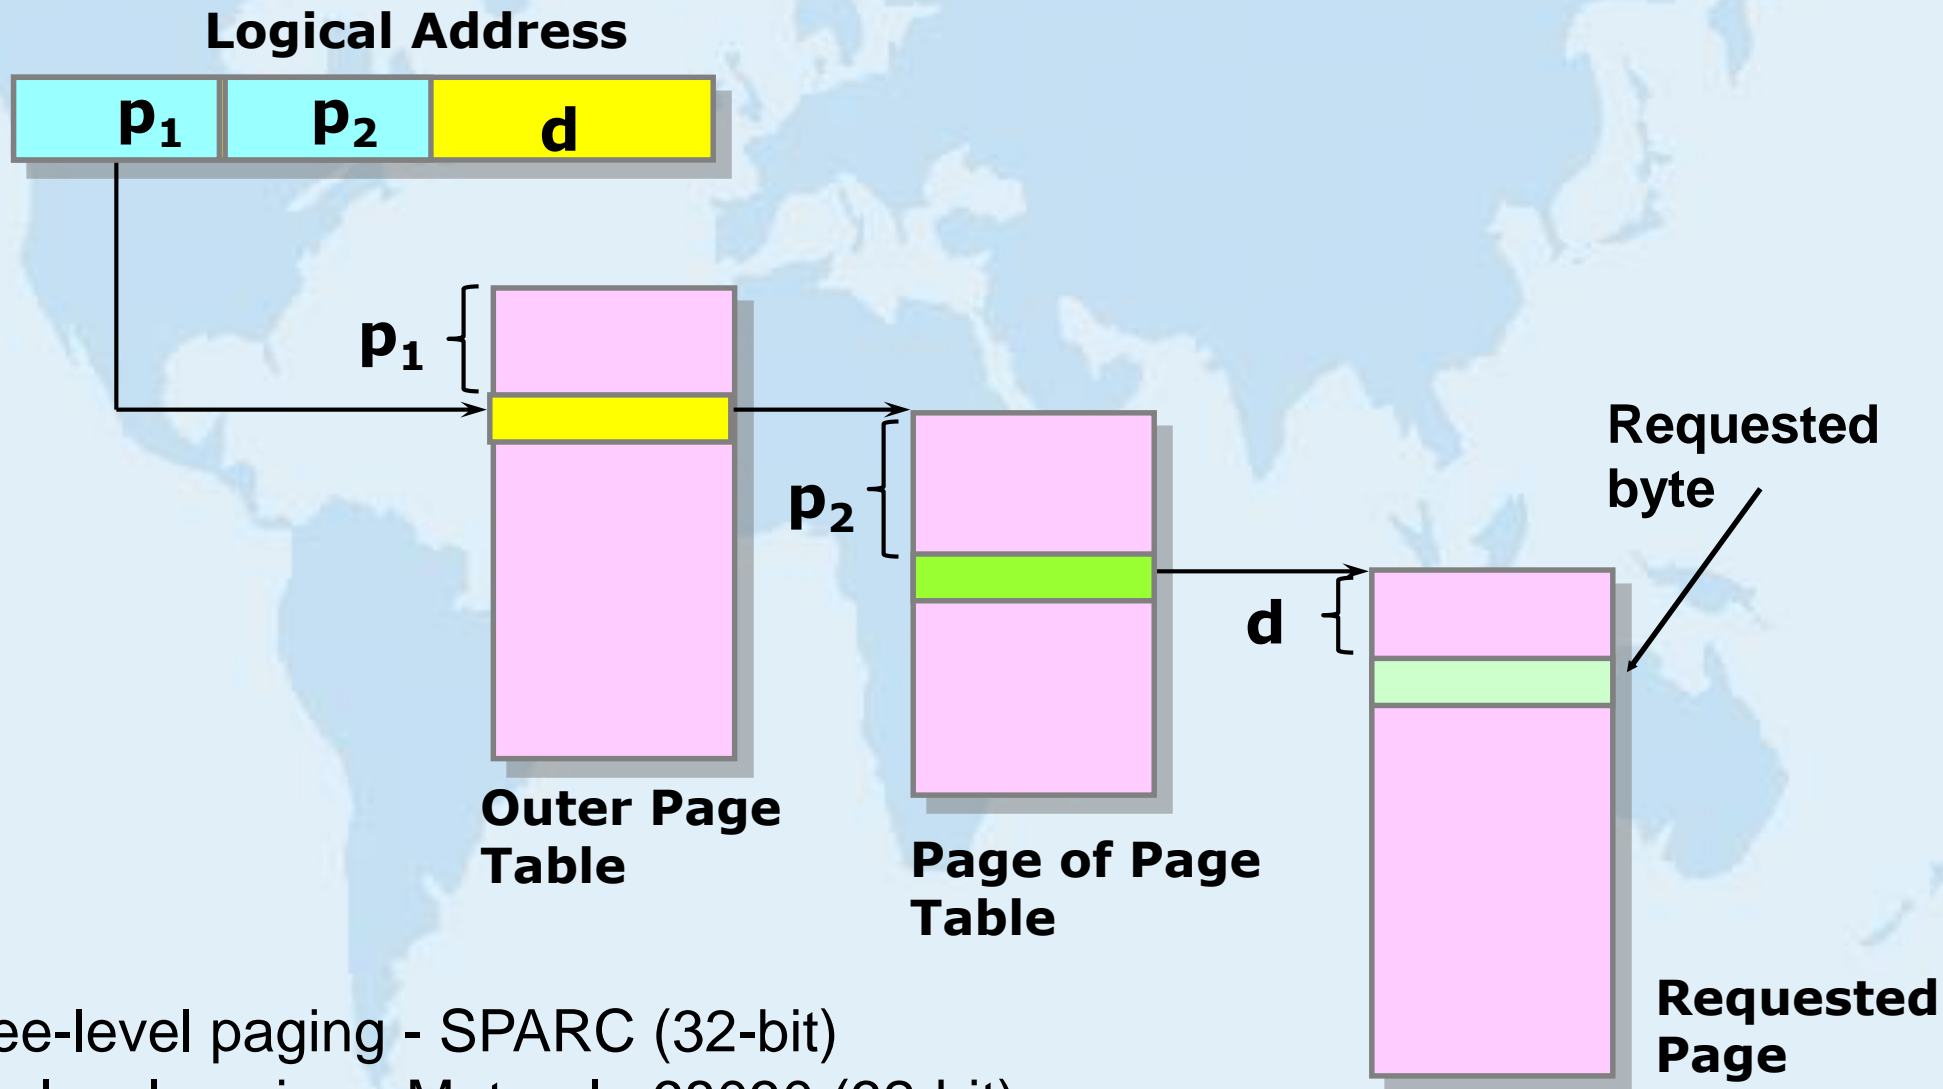
- **Example:** 32-bit machine with 4K page size.
- A logical address is divided into:
  - a **page number** consisting of 20 bits.
  - a **page offset** consisting of 12 bits.
- Since the Page Table is paged, the page number is further divided into:
  - a 10-bit page number  $p_1$ .
  - a 10 bit page offset  $p_2$ .
- Thus, a logical address is as follows:

Page Number		Page Offset
$p_1$	$p_2$	$d$
10	10	12

$p_1$ : index of the Outer Page Table.

$p_2$ : displacement within the page of the Outer-page Table.

# Two-level Paging



Three-level paging - SPARC (32-bit)

Four-level paging - Motorola 68030 (32-bit)



# Multilevel Paging Performance

- Since each level is stored as a separate table in memory, converting a logical address to a physical one may take up to **four** memory accesses.
- Even though time needed for one memory access is quintupled (five times as much), **caching** permits performance to remain reasonable.
- Cache hit ratio of 98% yields:

$$\text{EAT} = 0.98 \times 120 + 0.02 \times 520 = 128 \text{ ns.}$$

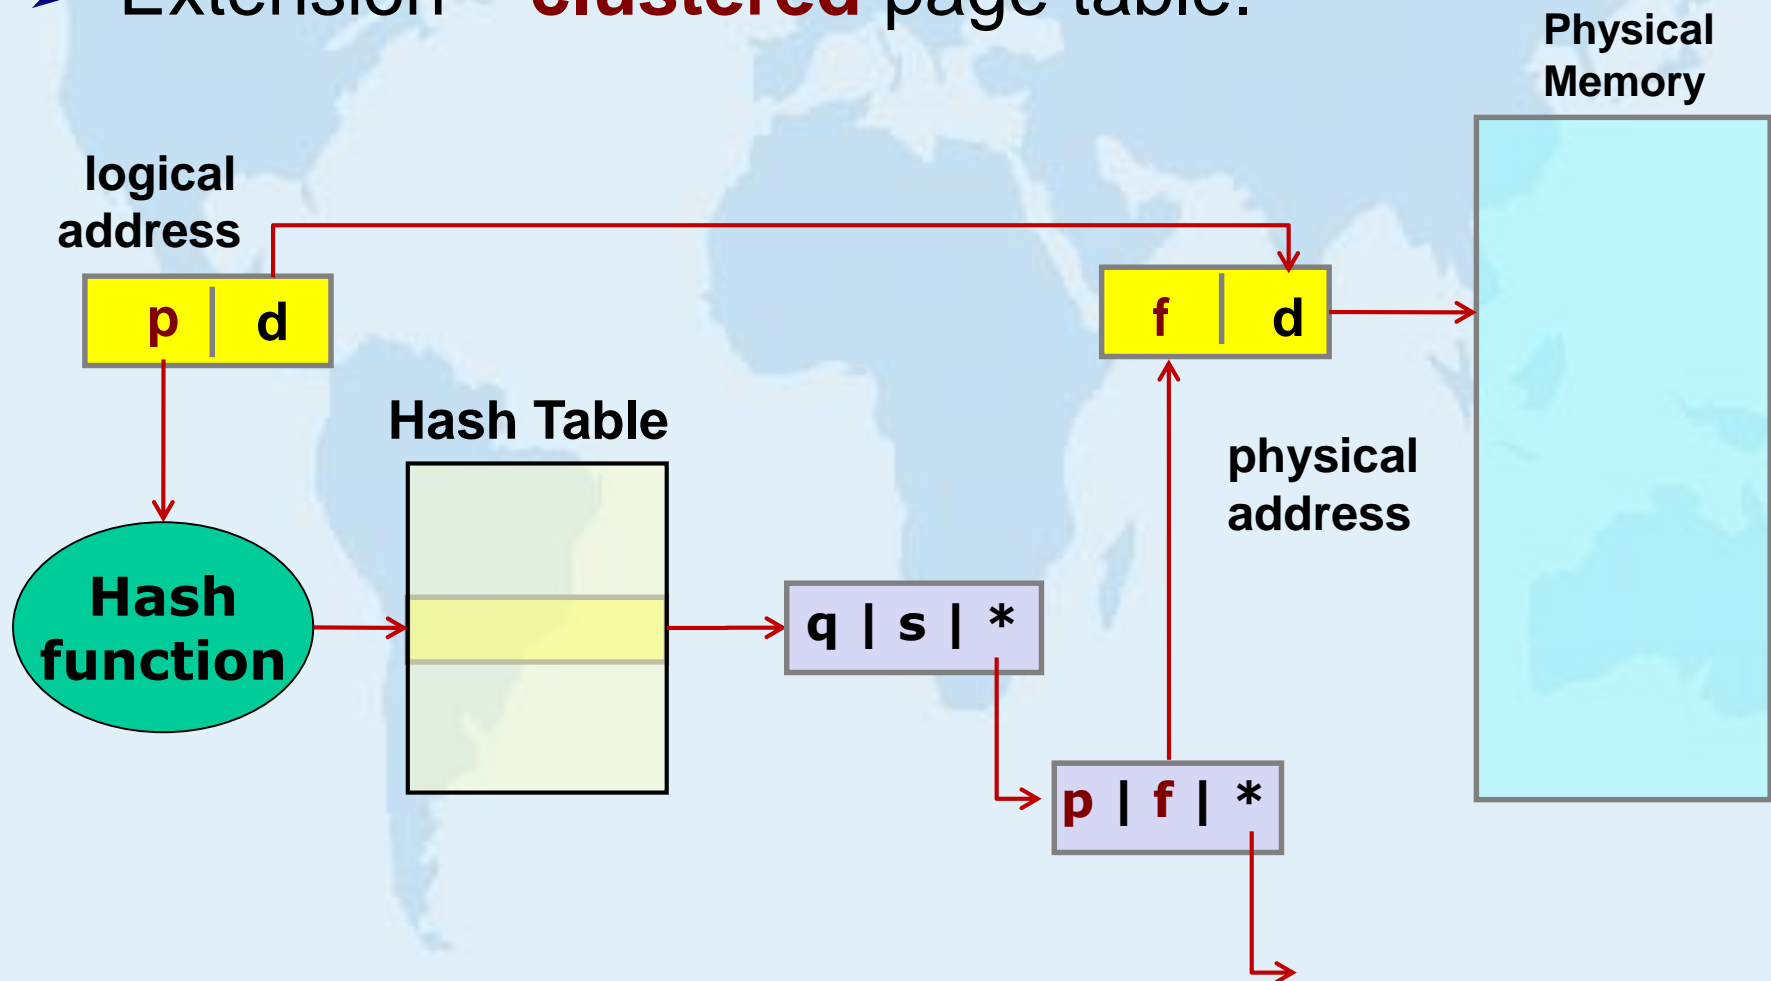
which is only a 28% slowdown in memory access time.

# Hashed Page Table

- For handling address spaces **larger** than 32 bits.
- Uses a **hash** table:
  - Hash value – virtual page number.
  - Each table entry - linked list of elements:
    - a) Virtual page number.
    - b) Value of mapped page frame.
    - c) Pointer to the next element.
- Algorithm:
  1. Virtual page number is **hashed** into the table.
  2. It is **compared** with the field a). If match, field b) is used.
  3. Else, **next** element is checked.

# Hashed Page Table

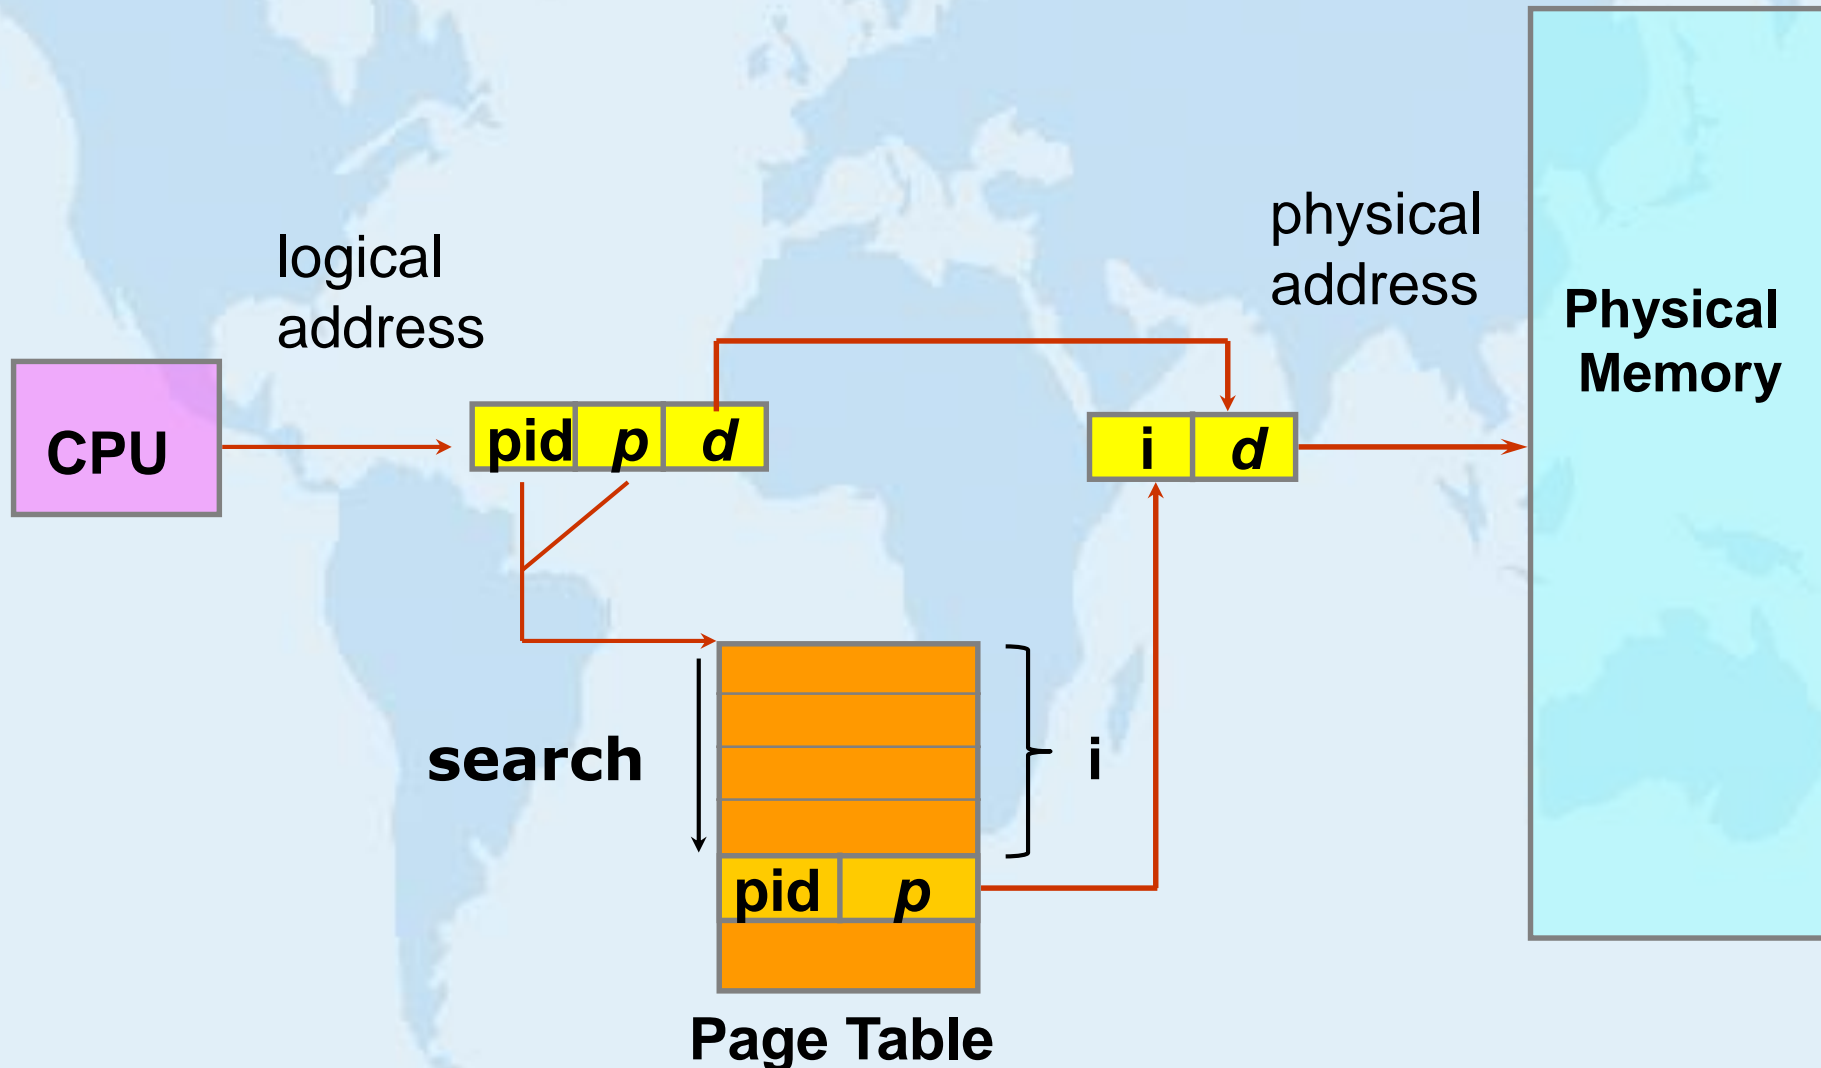
- Used for handling address spaces **larger** than 32 bits.
- Extension – **clustered** page table.



# Inverted Page Table

- **Inverted Page Table** - one entry for each real frame page of memory;
- **Inverted Page Table** - only one for all processes (before - each process has a own Page Table).
- Each table entry consists of the virtual address of the page stored in that real memory location, with information about the **process** that owns that page.
  - **Decreases** memory needed to store each Page Table.
  - **Increases** time needed to search the table.
- Example systems:
  - 64-bit UltraSparc, PowerPC


# Inverted Page Table



# Segmentation

- **Segmentation** - memory-management scheme that supports user view of memory.
- A program is a collection of segments. Each segment has a name and a length. A segment is a local unit such as:

- Main Program;
- Procedure;
- Function;
- Local / Global Variables;
- Common Block;
- Stack;
- Symbol Table, Arrays.



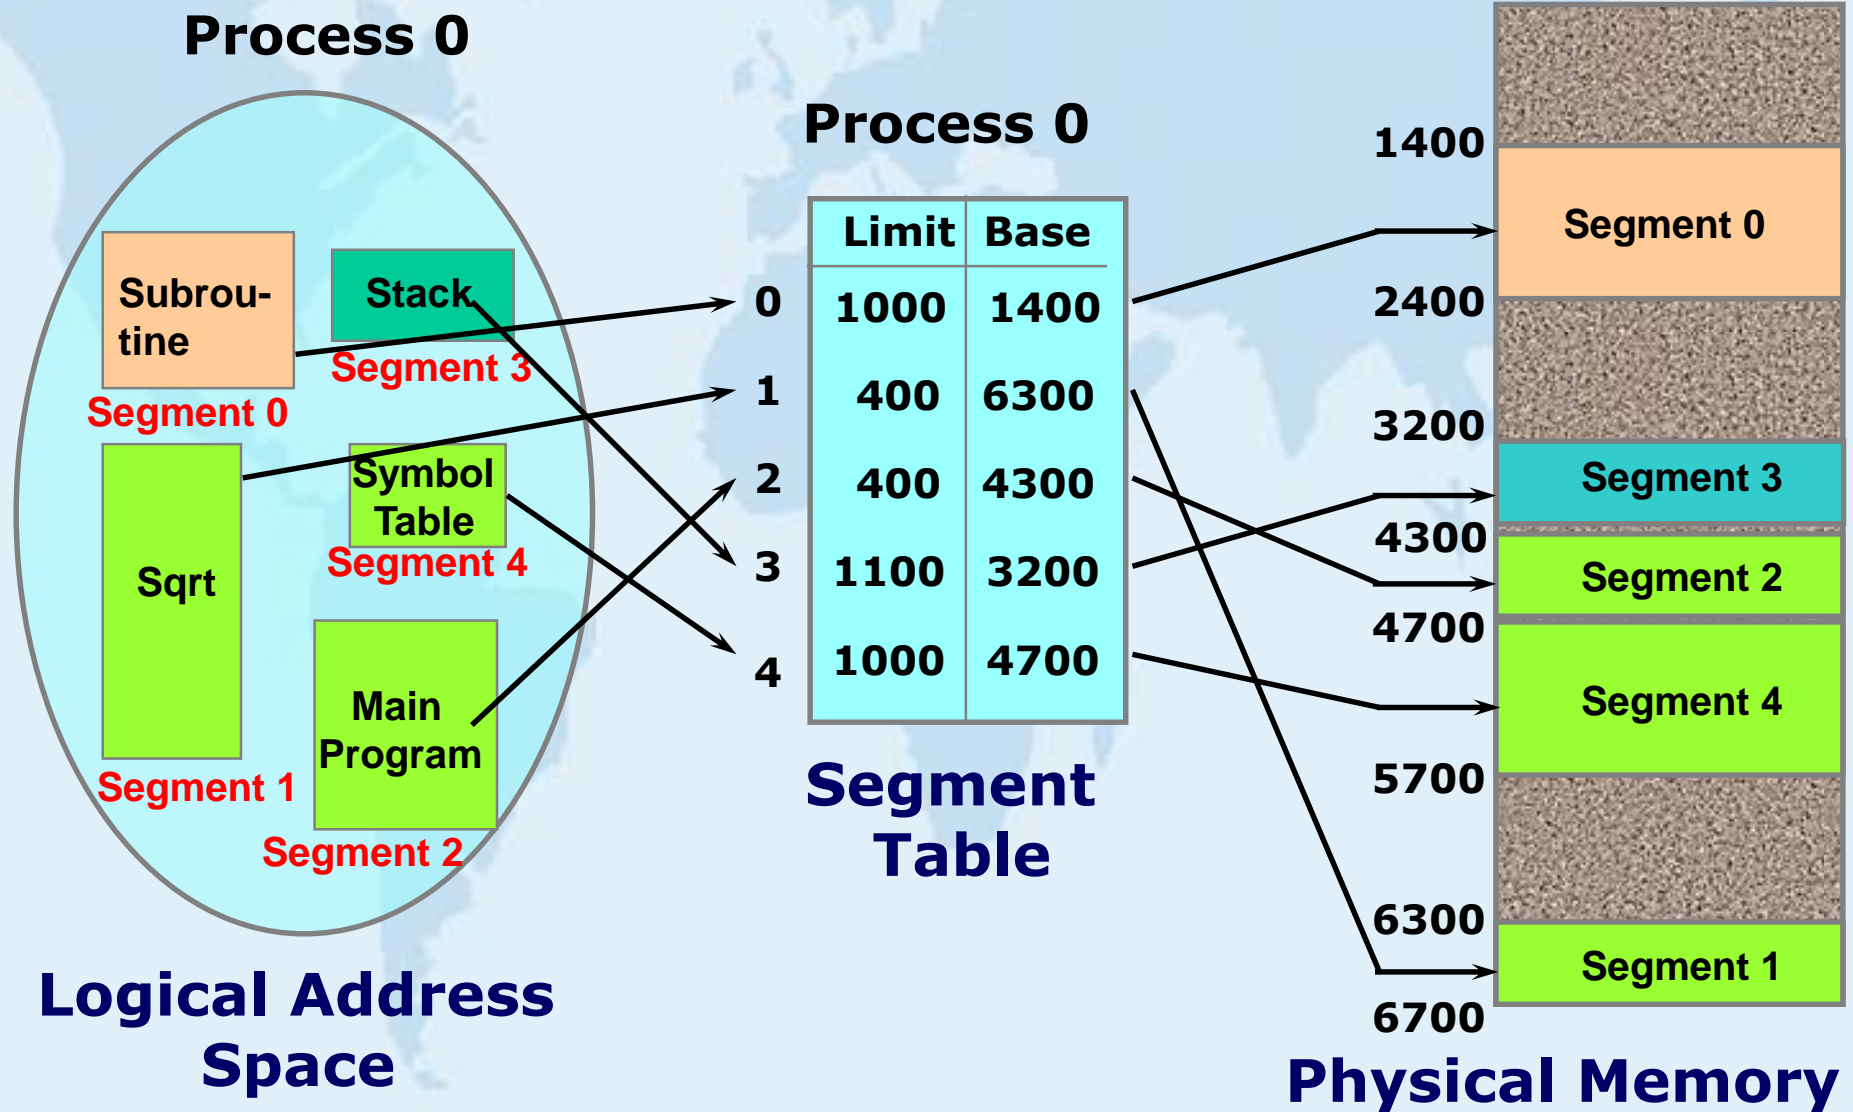


# Segmentation

- Logical address consists of following tuple:  
**< segment-number, offset >**
- **Segment Table** - maps two-dimensional user-defined addresses into one-dimensional physical addresses;
- Each entry of the **Segment Table** has:
  - **base** - contains the starting physical address where the segments reside in memory.
  - **limit** - specifies the length of the segment.

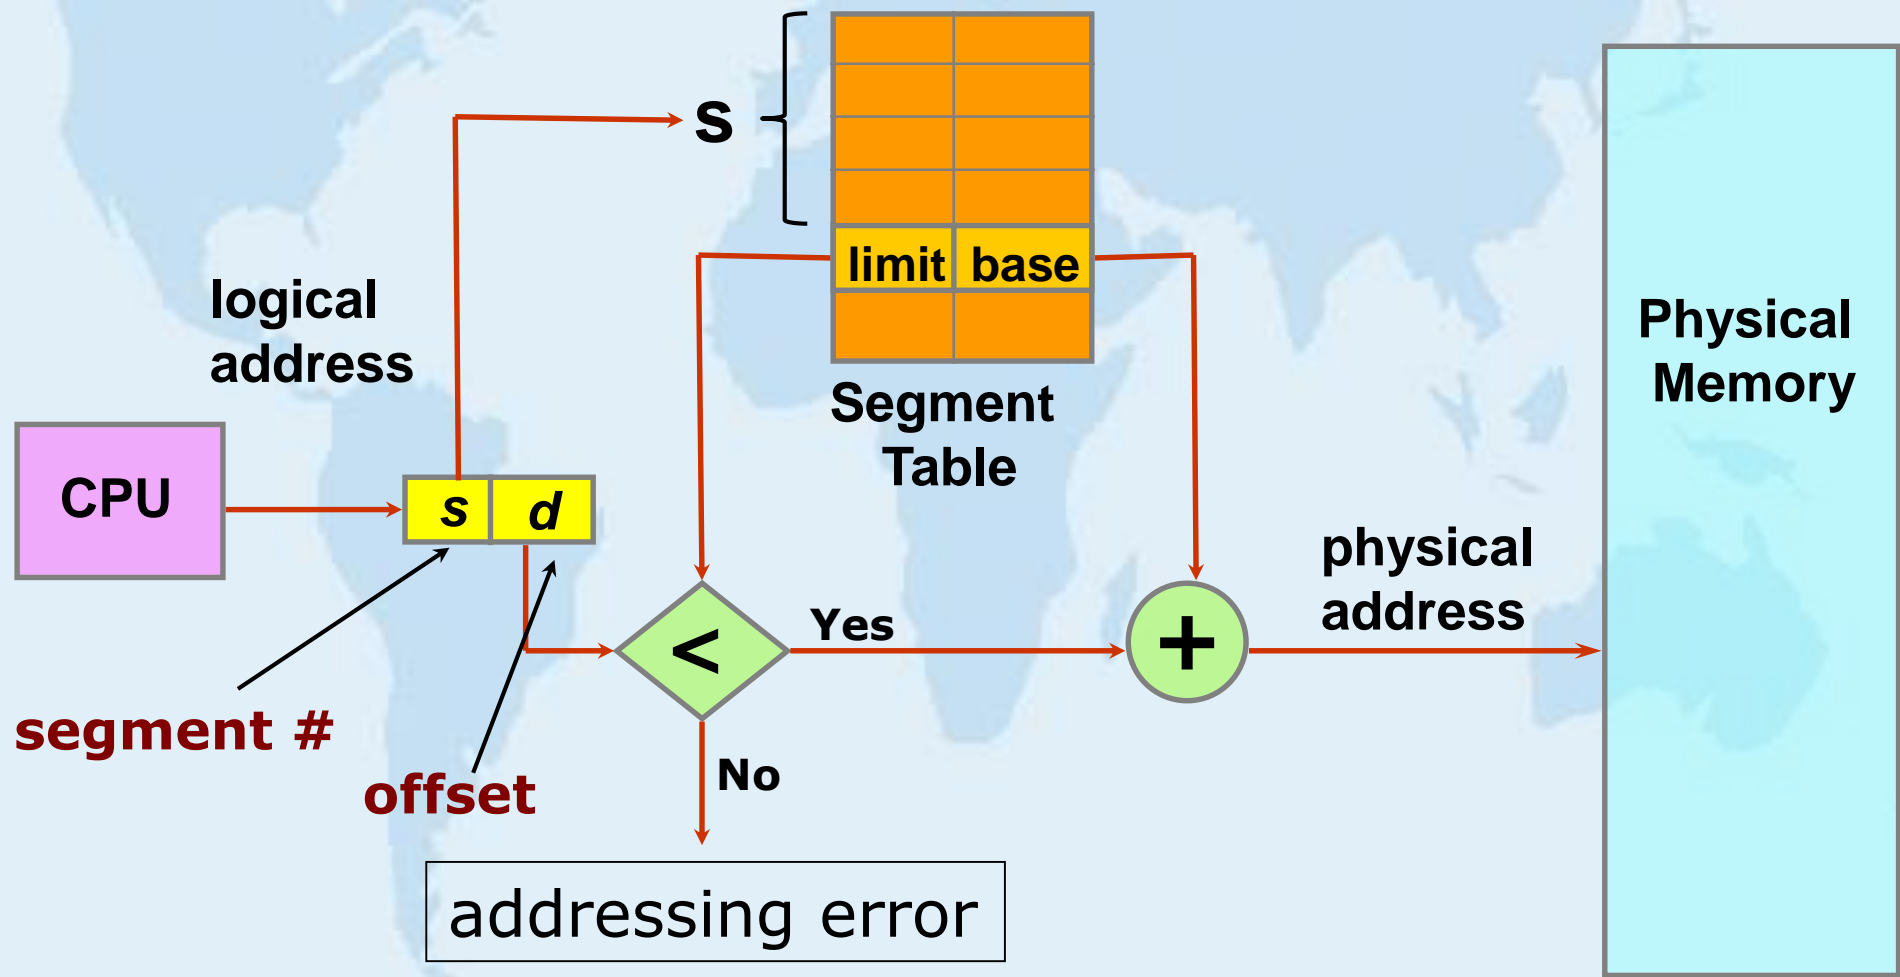


# Segmentation



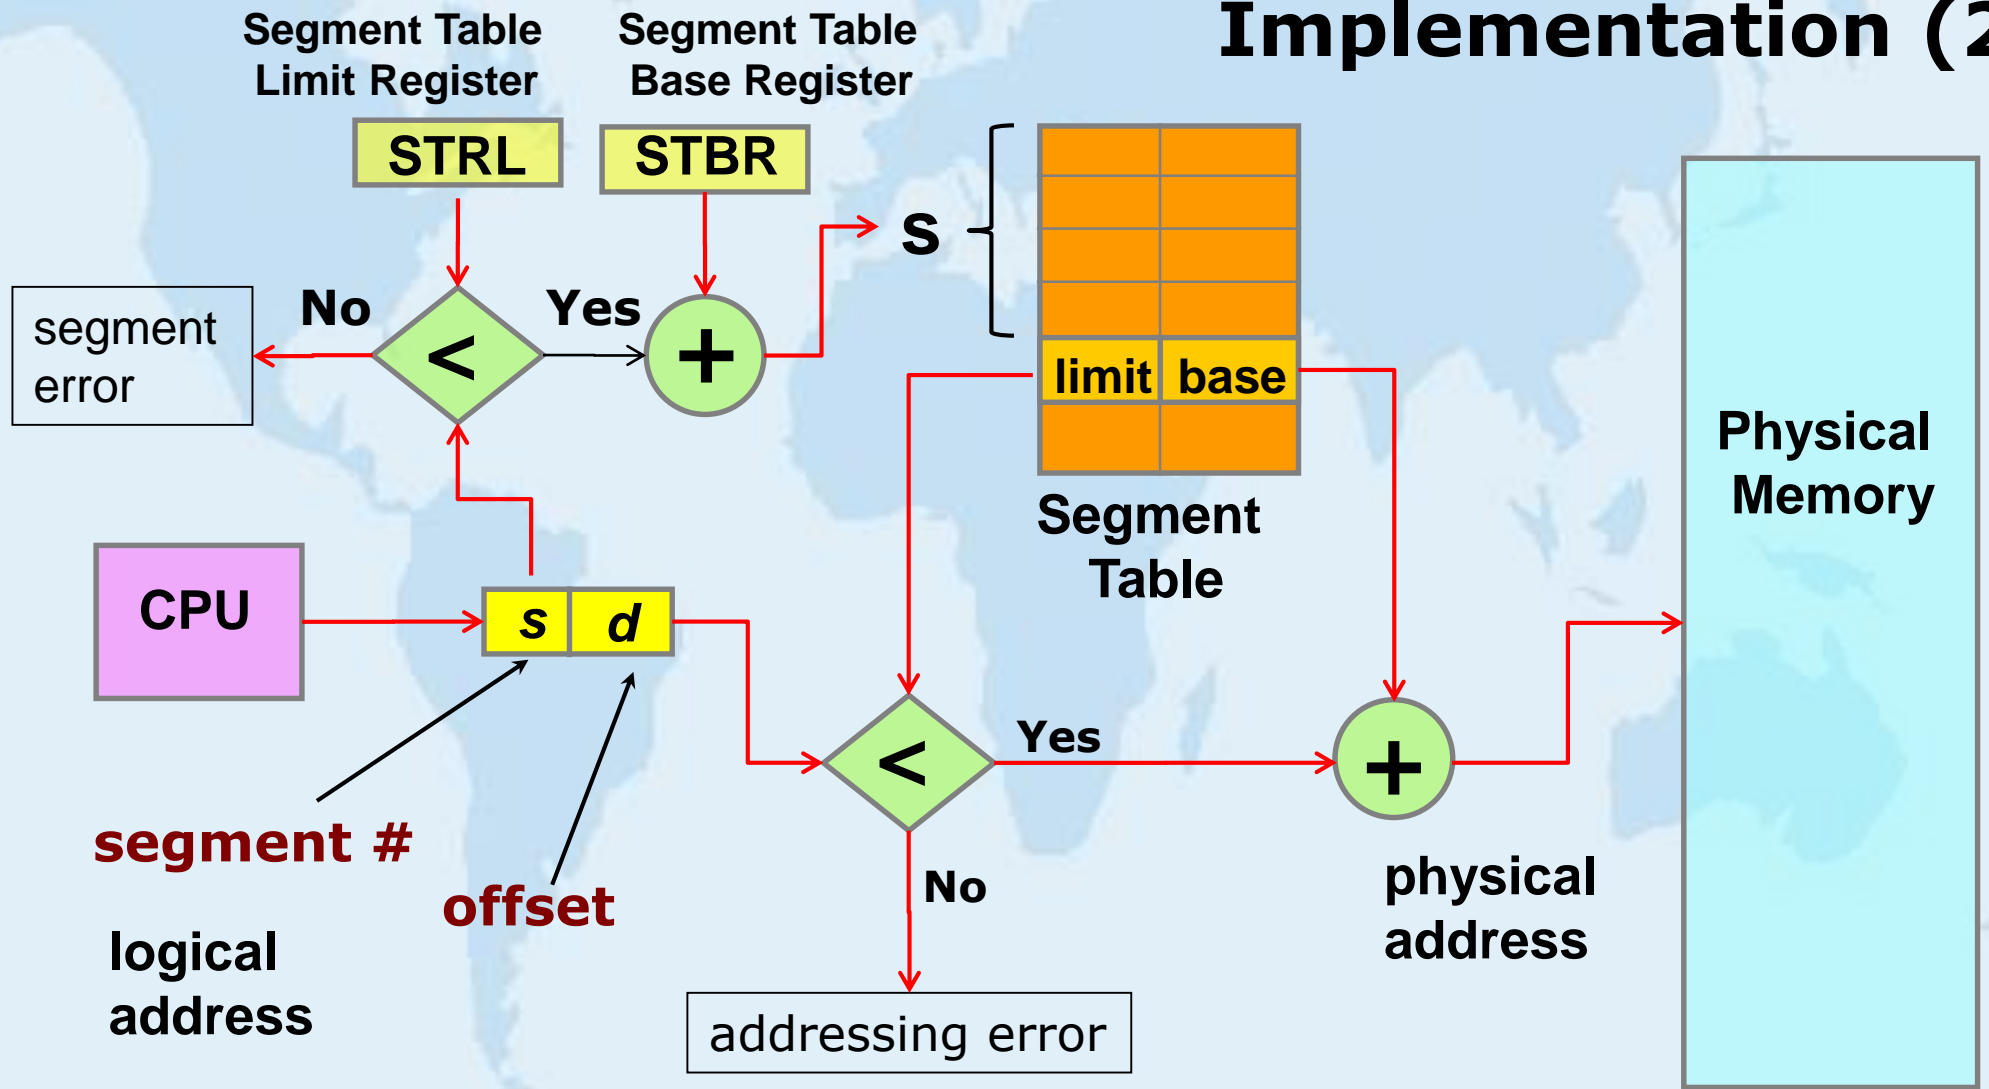
# Segmentation

## Implementation (1)



# Segmentation

## Implementation (2)




# Segmentation

- **Protection** and **Sharing** of data can be implemented naturally at the segment level.
- **Protection:** segment - semantically defined portion of program;
  - instructions ➡ execute-only or read-only;
  - data ➡ read- or write-only);
  - protection bits associated with segments.
- **Sharing:** segments are shared when entries in the Segment Tables of two different processes point to the same physical locations.


# Segmentation - Sharing

Process # 0



Logical Memory Process # 0

Process # 1



Logical Memory Process # 1


Segment Table  
Process # 0

	Limit	Base
0	25286	43062
1	4425	68348

Segment Table  
Process # 1

	Limit	Base
0	25286	43062
1	8850	90003

Physical Memory



# Segmentation-Issues

- Since segments vary in length, memory allocation is **dynamic storage-allocation** problem (best-fit or first-fit).
- Segmentation may cause **external fragmentation**, when all blocks of free memory (holes) are too small to accommodate a segment.
- But we can compact memory whenever we want. Average segment size?
  - **One extreme:** each process to be one segment.
  - **The other extreme:** every byte put in its own segment (doubling memory use!).



# Segmentation-Issues

- The next logical step - fixed-sized, small segments - is **paging**.
- It is possible to combine **Segmentation** and **Paging** to improve on each.
- This combination is best illustrated by two different architectures: MULTICS system and Intel 386.



# Segmentation with Paging

- **MULTICS Operating System**
  - Segmentation with Paging.
  - Dynamic linking.
  - Ring protection.
  - Most code was written in PL/1.
- **Hardware**
  - GE 635 was modified into GE 645.
  - Word machine; 1 word = 36 bits.
  - Max. # of segments = 256 K.
  - Segment size = 64 K words.
  - Page size = 1 K words.



**That is all for today!**